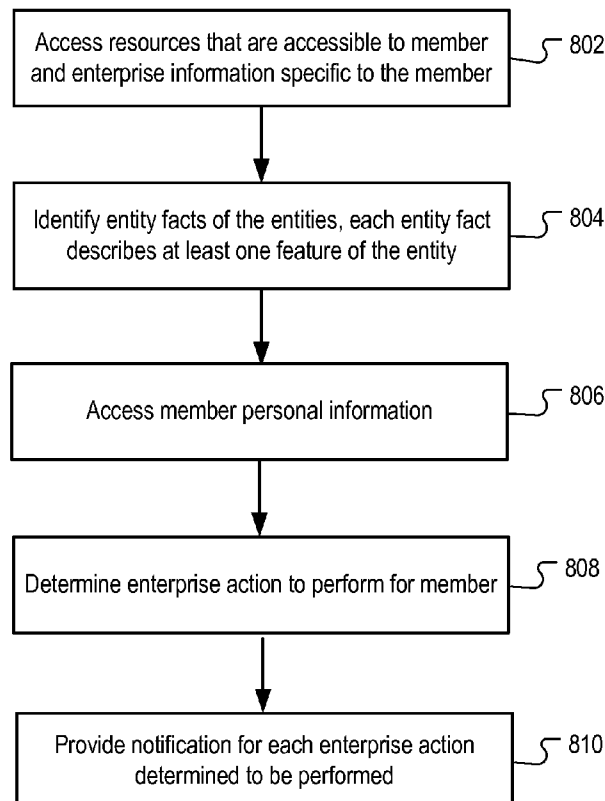




US 20160352743A1

(19) **United States**(12) **Patent Application Publication**  
**VerWeyst et al.**(10) **Pub. No.: US 2016/0352743 A1**(43) **Pub. Date: Dec. 1, 2016**(54) **ACCESS CONTROL FOR ENTERPRISE  
KNOWLEDGE**(71) Applicant: **Google Inc.**, Mountain View, CA (US)(72) Inventors: **Brent VerWeyst**, San Francisco, CA  
(US); **Martin James Cochran**, San  
Francisco, CA (US); **Muthian  
Sivathanu**, Chennai (IN)(21) Appl. No.: **14/723,760**(22) Filed: **May 28, 2015****Publication Classification**(51) **Int. Cl.**  
**H04L 29/06** (2006.01)  
**G06F 21/62** (2006.01)  
**G06F 17/30** (2006.01)(52) **U.S. Cl.**  
CPC ..... **H04L 63/101** (2013.01); **G06F 17/30011**  
(2013.01); **G06F 21/6218** (2013.01)(57) **ABSTRACT**

Methods, systems, and apparatus, including computer programs encoded on a computer storage medium for access control for enterprise information. In one aspect, a method includes receiving resources of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members, and the resources including entities related to the enterprise and relationships; identifying entity facts of the entities from the resources; determining, for each entity fact, an entity fact access control list; storing data describing the entities, entity facts and the respective entity fact access control lists, wherein each entity fact is associated with its corresponding entity fact access control list; and providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts according to the respective entity fact access control lists.

800  





FIG. 1




FIG. 2

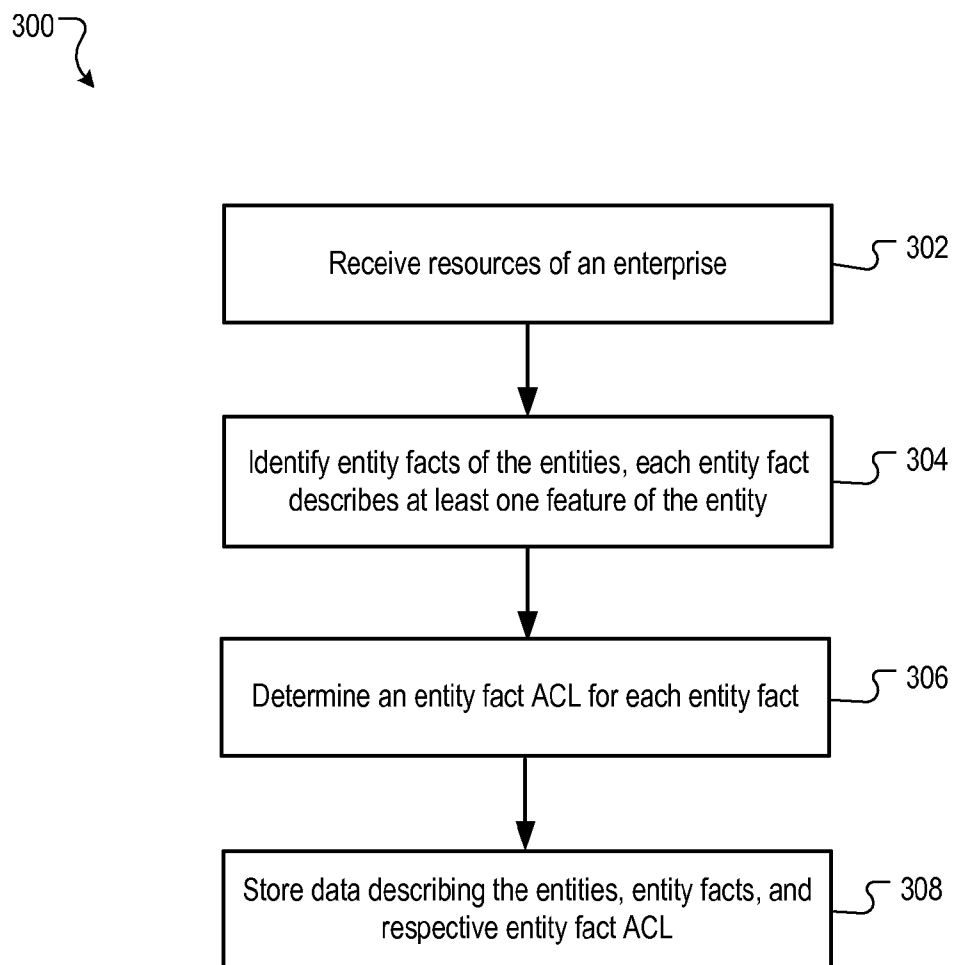


FIG. 3

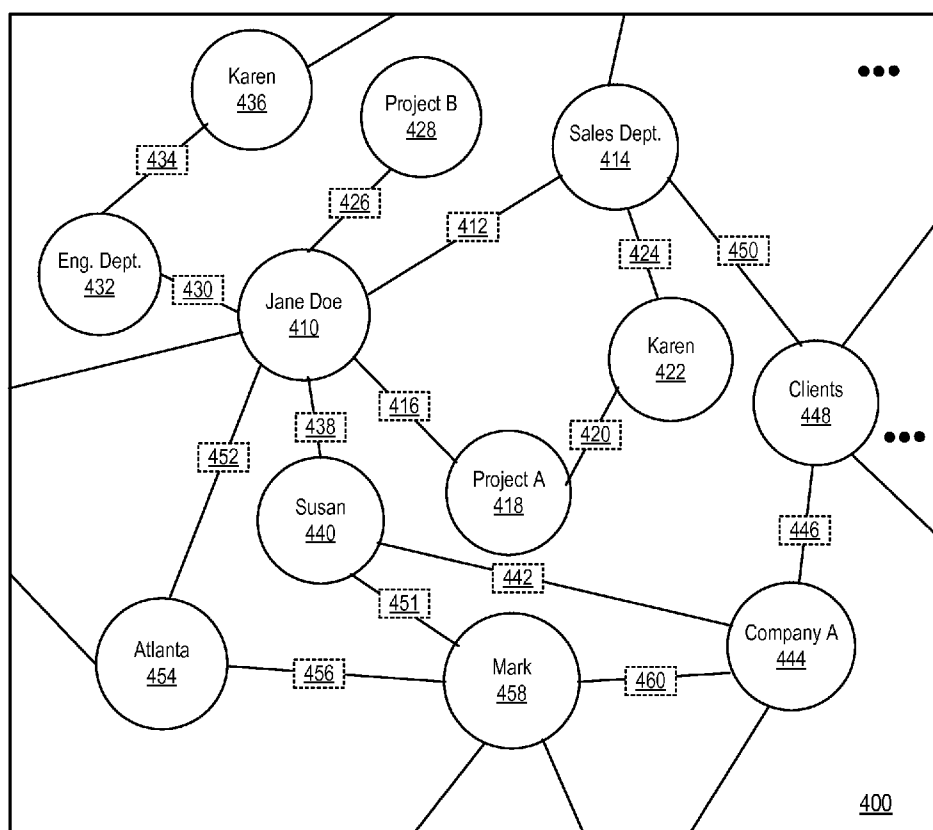


FIG. 4

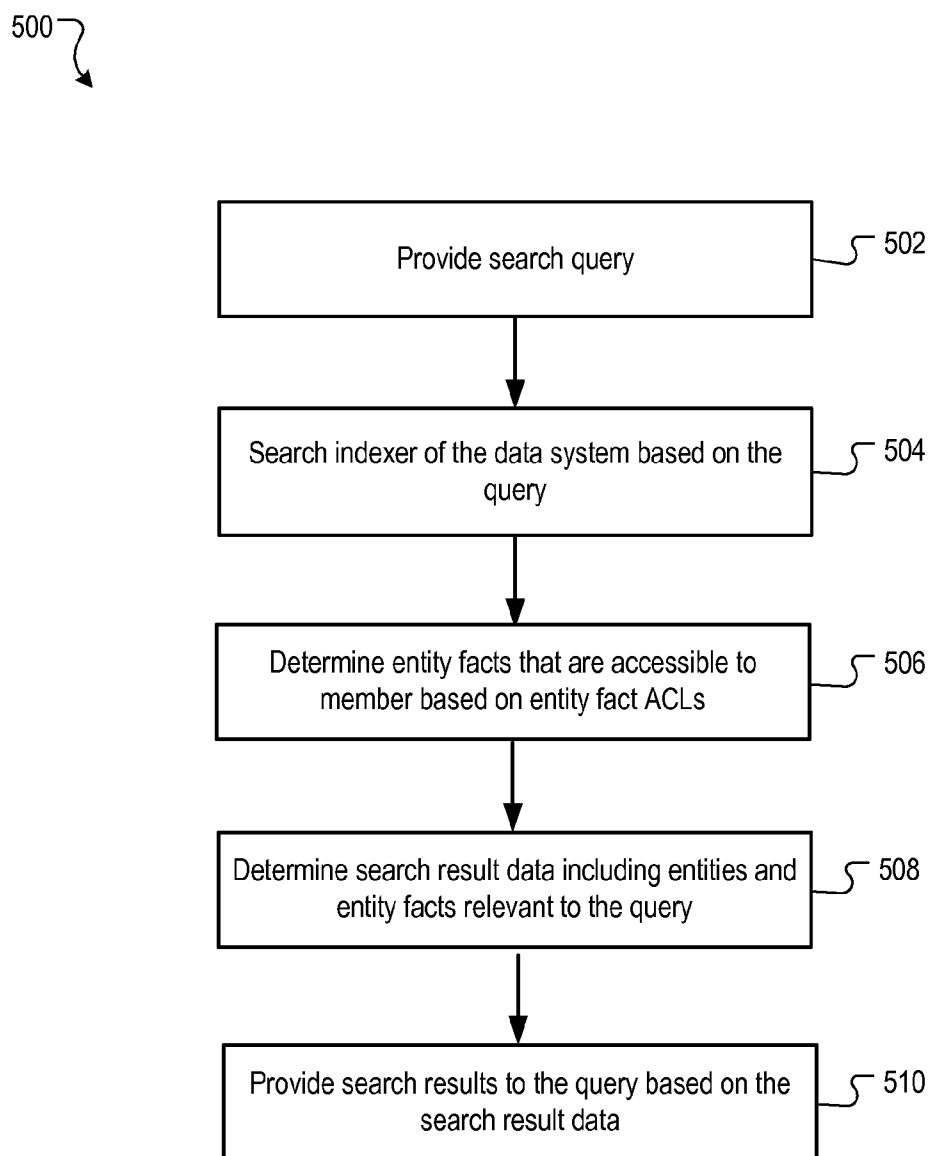


FIG. 5

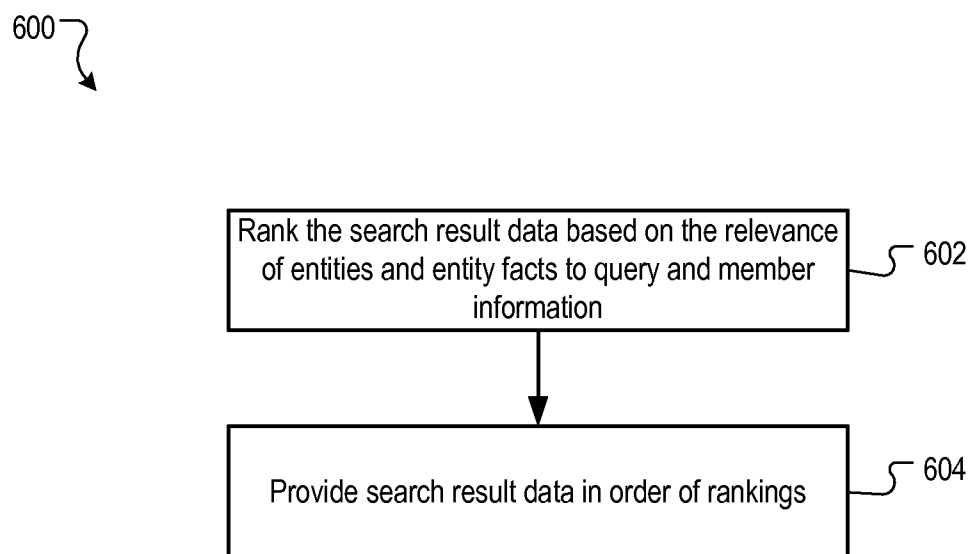
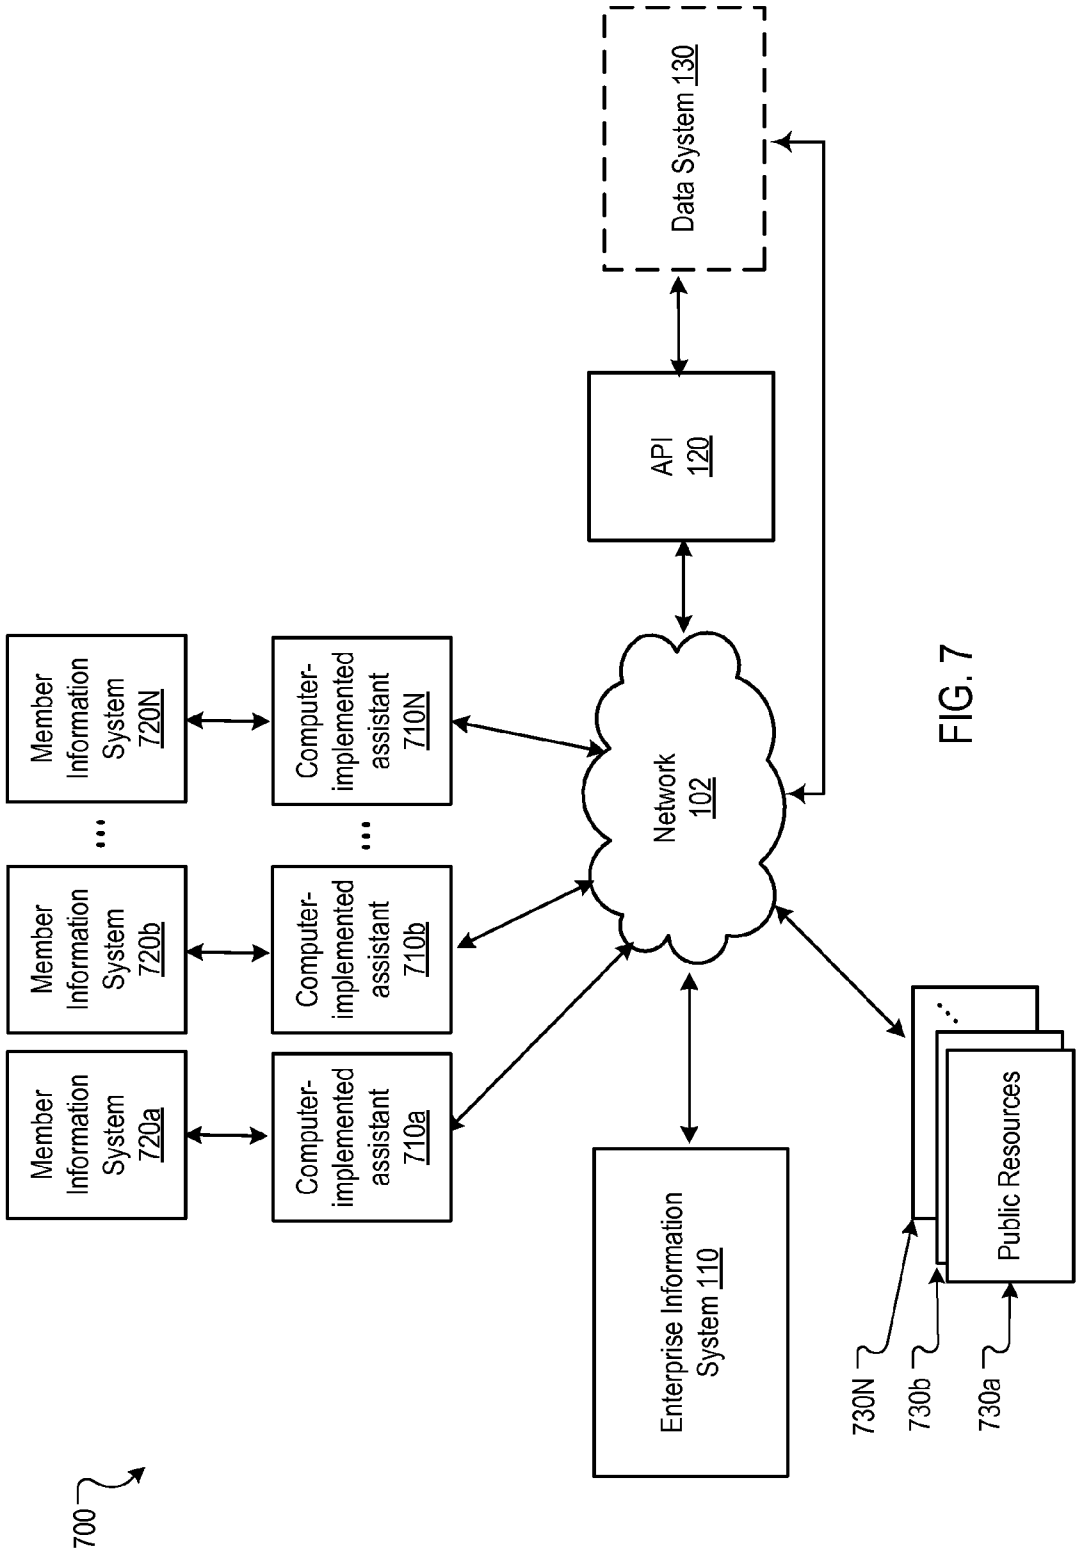


FIG. 6





800

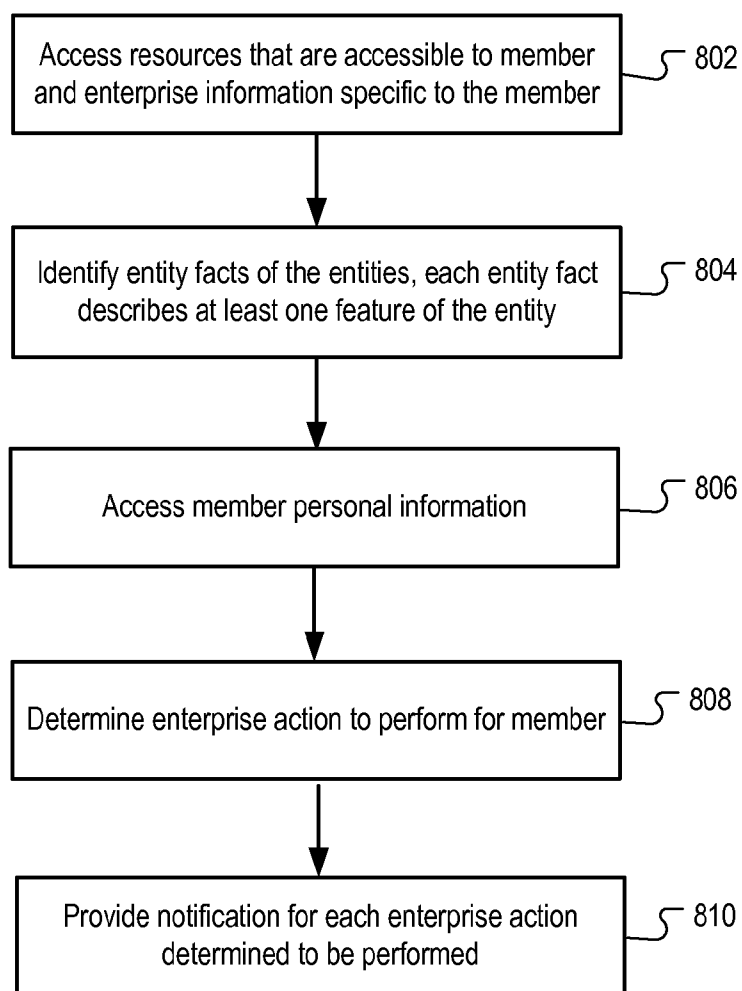


FIG. 8

## ACCESS CONTROL FOR ENTERPRISE KNOWLEDGE

### BACKGROUND

[0001] This specification relates to access control for enterprise information, personal assistance based on enterprise information and personal information, and searches associated with the enterprise information.

[0002] Enterprises can have access to a large number and wide variety of resources (e.g., documents). Each of the resources may have access control lists that are associated with the resource to dictate who can view and/or alter the resource. Members of an enterprise may have difficulty in finding and searching for necessary information because the resources are not connected or associated with one another in any meaningful way. Furthermore, and for similar reasons, electronic personal assistants typically cannot leverage enterprise information in meaningful ways to satisfy users' informational needs.

### SUMMARY

[0003] This specification describes technologies relating to access control for enterprise information, personal assistance based on enterprise information and personal information, and searches associated with the enterprise information.

[0004] In general, one innovative aspect of the subject matter described in this specification can be embodied in a method that includes the actions of receiving resources of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members of the enterprise, and the resources including data describing entities related to the enterprise and relationships among the entities. Entity facts may be identified from the entities from the resources of the enterprise, each entity fact describing at least one feature of the entity, wherein the features include corresponding relationships between the entities. For each entity fact, a determination of an entity fact access control list may be made from the access control list of each resource, an entity fact access control list, and the data describing the entities, entity facts and the respective entity fact access control lists may be stored in a searchable index, wherein each entity fact is associated with its corresponding entity fact access control list. Also, the method may include providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists.

[0005] Particular implementations of the subject matter described in this specification can be implemented so as to realize one or more of the following advantages. In some enterprise structures, it can be beneficial to extract information from the resources and compile that information while still respecting the access control list of the resource with the extracted information, as this will cause the extracted information to be easier to obtain by members of the enterprise while also protecting the access control lists associated with the resources. Also, it may be beneficial to have a computer-implemented person assistant to perform actions, where the computer-implemented personal assistant uses information

from each member and extracted information from the resources while respecting the access control lists of the resources.

[0006] The details of one or more embodiments of the subject matter described in this specification are set forth in the accompanying drawings and the description below. Other features, aspects, and advantages of the subject matter will become apparent from the description, the drawings, and the claims.

### BRIEF DESCRIPTION OF THE DRAWINGS

[0007] FIG. 1 is a block diagram of an example environment in which enterprise data is integrated.

[0008] FIG. 2 is a block diagram of a data system.

[0009] FIG. 3 is a flow diagram of an example process for providing access privileges to entity facts.

[0010] FIG. 4 is a block diagram of a portion of an example enterprise knowledge graph.

[0011] FIG. 5 is a flow chart of an example process for providing search results to an enterprise member.

[0012] FIG. 6 is a flow chart of an example process for ranking and providing the search results to an enterprise member.

[0013] FIG. 7 is a block diagram that incorporates computer-implemented assistants into the example environment structure of FIG. 1.

[0014] FIG. 8 is a flow chart of an example process for providing enterprise actions by a computer-implemented assistant.

### DETAILED DESCRIPTION

[0015] Overview

[0016] FIG. 1 is a block diagram of an example environment 100 in which enterprise data is integrated. The example environment 100 includes a network 102, e.g., a local area network (LAN), wide area network (WAN), the Internet, or a combination of them, connects enterprise information system 110, application programming interface 120, and data system 130. The network 102 can be accessed over a wired and/or a wireless communications link. For example, mobile computing devices, such as smartphones can utilize a cellular network to access the network.

[0017] Enterprise information system 110 may include one or more information sources 112. As seen in FIG. 1, information sources 112 are identified as 112a, 112b, . . . , 112N, which include one or more sub-system, application, program, or database, among others. For example, information sources 112 may be an internal document database for an enterprise, an enterprise intranet, an enterprise email application, or any third party information source that the enterprise information system 110 is connected to (e.g., internet, software application). Each of the information sources 112 include resources 115, which in FIG. 1, may be identified as 115a1, 115a2 . . . 115an . . . 115n1 . . . 115NN. Resources 115 may be different for each information source 112. For example, resources 115 may be a document stored within the enterprise information system 110, intranet URL, email message, internet URL, or any similar type of informational content (e.g., information stored on third party application that the enterprise information system 110 has access to).

[0018] As used in this document, an "enterprise" is public or private legal entity to which members belong and for

which the members undertake projects, tasks, and other actions on behalf of the enterprise. An example enterprise may be a privately held company, a publically traded company, and the like. “Enterprise information” is at least information that is authored, owned, curated or otherwise controlled, either partially or exclusively, by the enterprise.

**[0019]** Each resource **115** has an associated access control list (ACL) that indicates which members of the enterprise are allowed to access and/or modify that particular resource. The access control list for each resource may indicate access privileges, such as an owner (or creator) of the resource, read access, write access, and delete access, among others. The access control list for each resource may define access and/or modification privileges to individual members and/or one or more groups of members (e.g., sales department of an enterprise). Each resource **115** may also allow for versioning of the particular resource (e.g., version **1** of resource **115a1** and version **2** of resource **115a1**), and the access control list for each version of the particular resource may be different from one another. Also, a particular resource may be created or defined as a “parent” or “child” of another resource. The access control list for any child resource will inherit the access control list for its defined parent resource; however, in some implementations, this is only a default access control list and the access control list for the child resource and/or parent resource may be modified without affecting the other. In variations of this implementation, a child resource may have additional access control list requirements beyond those inherited from the parent resource.

**[0020]** Resources **115** may be provided over network **102** to a data system **130** via an application programming interface (API) **120**. API **120** may be designed by the enterprise, the operator of data system **130**, or a third party (e.g., developed by a third party for a specific application or a more generally applicable API). API **120** may accept different access control list types that are used or presented in information sources **112** that include more or fewer different definitions or fields (e.g., owner, read or write access, parent, child, versions, etc.). Additionally, API **120** allows for multiple identifications for each member of the enterprise. For example, if a member has an identification of “Jane Doe” for information source **112a** and has identification “JaneDoe@enterprise.com” for information source **112b**, and the enterprise information system **110** has provided a description that both of these identifications relate to the same member, the API **120** can link both of these identifications with that particular member. Additionally, API **120** enables members to be classified in one or more group, and if the group is updated, added, or removed on at the enterprise information system **110**, then the API **120** may sync with the enterprise information system **110** to reflect the changes at the API **120**.

**[0021]** The resources **115** are received through the API **120** by the data system **130**. A database schema implemented by data system **130** provides for the data system **130** to understand the information provided by enterprise information system **110** via the API **120** in order to integrate and structure the information. Further, the API **120** may tag the information associated with resources **115** according to the predefined schema in order for the data system **130** to integrate the information. In the current implementation, data system **130** includes an integration system **132**, a data storage **134**, and an indexer **136**. Integration system **132** may be a database, server, module, application, or combination

thereof, among others. Information from the integration system **132** is stored in data storage **134**, and indexer **136** indexes the information stored in data storage **134** for the information to be searched by the members of the enterprise, which will be described below.

**[0022]** Entity Fact Access Control

**[0023]** FIG. 2 is a block diagram of data system **130**. In some implementations, integration system **132** receives a plurality of resources **115** that each have an access control list specifying access privileges to the resource **115** for one or more members of the enterprise. Each resource includes data and/or information describing entities **210**, which may be identified as entities **210a**, **210b** . . . **210N** in FIG. 2. The entities **210** are related to the enterprise and relationships between the entities. Entities **210** of the enterprise may be, for example, people (or members), positions, projects, departments, sales information, performance evaluations, or any other information that is in the resources **115** of the enterprise. More generally, in some implementations, entities are topics of discourse, concepts or things that each have a distinct and independent existence and that can be referred to by a text fragment, e.g., a term or phrase, and are distinguishable from one another, e.g., based on context. An entity may thus be a physical or conceptual realization having a factual definition and attributes.

**[0024]** Additionally, integration system **132** derives entity facts **220**, which may be identified as entity facts **220a**, **220b**, . . . **220N** in FIG. 2 from the entities **210** of the resources **115**. Each entity fact **220** describes at least one feature of an entity **210**, where the features describe corresponding relationships between the entities **210**. For example, a resource **115** (e.g., a document) of the enterprise may identify the names and positions of those in management of the enterprise where “Jane Doe” may be in management of the enterprise in the position of “Vice-President of Sales.” In the current example, “Jane Doe” and “Vice-President of Sales” are entities, and an entity fact **220** may be defined in the integration system **132** that creates a relationship between the entities, where the entity relation may be expressed as a factual relation “Jane Doe”—“Vice-President of Sales.” A variety of techniques may be used to identify entity facts. These include linguistic parsers that can be issued to determine, e.g., nouns and related attributes; database parsers that can be used to determine keyed relations between date fields; and so on.

**[0025]** Further, integration system **132** determines an entity fact access control list **230**, which may be identified as **230a**, **230b**, . . . **230N** in FIG. 2. An entity fact access control list **230** is determined for each entity fact **220**, which is determined from the access control list of each resource **115**. In some implementations, the access control list of a resource is inherited by each entity fact derived from the resource. For example, resource **115a1** may be a document with salary information of the members of the enterprise. Resource **115a1** may include “Jane Doe” as entity **210a** and “\$50,000” as entity **210b** in a separate column (indicating Jane Doe’s salary). As salary information is personal and sensitive information, the access control list for resource **115a1** may only include the Director of Human Resources, for example, “Mary.” Resource **115b1** may be an intranet page that breaks down the salary amount of members by percentages, and may include “Jane Doe” as entity **211a** and “Top 5%” as entity **211b** (indicating Jane Doe’s income percentage). The access control list for this information may

provide access to all members of the enterprise (or a larger group of the members than resource **115a1**). Entity facts may then be identified, and in the present example, may be “Jane Doe” and “\$50,000” salary as entity fact **220a**, and “Jane Doe” and “Top 5%” as entity fact **220b**, which are derived from entities **210a**, **210b**, **211a**, and **211b** of resources **115a1** and **115b1**.

[0026] Entity fact access control lists **230** may then be determined based on the resource from which the respective entity fact was derived. For example, entity fact access control list **230a** for entity fact **220a** will only include Mary because resource **115a1**, from which entity fact **220a** was derived, included Mary in the access control list. Entity fact access control list **230b** for entity fact **220b** will include all members of the enterprise because resource **115b1** included all members of the enterprise in the access control list. Based on the current example, Mary would have access to entity fact **220a** and entity fact **220b** because she is included in the access control list for both resource **115a1** and **115b1**; however, other members, for example, “Alvin,” a sales representative, would not have access to **220a** because he is not part of the access control list for the underlying resource **115a1**, but he would have access to entity fact **220b** because he has access privileges to resource **115b1** based on the access control list associated with the document from which the fact **220b** was derived.

[0027] The information sent to and obtained by the integration system **132** is stored in data storage **134**. Data storage **134** stores, for example, data describing the entities **210**, entity facts **220** and the respective entity fact access control lists **230**, where each entity fact **220** is associated with its corresponding entity fact access control list **230**. Indexer **136** then indexes the information in data storage **134**. The members of the enterprise may then use a search engine **138** to, for example, search data storage **134** to access entity facts **220** of the resources **115** of the enterprise. Each member of the enterprise is provided access privileges to the data describing the entities **210** and the entity facts **220** according to the respective entity fact access control list **230**.

[0028] Any appropriate search process can be used to realize the search engine **138**. The search engine **138** can use, for example, one or more of an information retrieval (IR) score algorithm, an authority score algorithm, or a combination of relevance and authority algorithms.

[0029] FIG. 3 is a flow chart of an example process **300** for providing access privileges to entity facts **220** based on respective entity fact access control lists **230** obtained from underlying resources **115** of the enterprise. The process **300** can, for example, be implemented by the data system **130**. In some implementations, the operations of the example process **300** can be implemented as instructions stored on a non-transitory computer readable medium, where the instructions cause a data processing apparatus to perform operations of the example process **300**.

[0030] Resources **115** of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members of the enterprise, and the resources including data describing entities **210** related to the enterprise and relationships among the entities **210** are received (**302**). Each information source **112** of the enterprise information system **110** includes resources **115**, and the resources **115** may be provided over network **102** to a data system **130** via the API **120**. Each resource **115** includes an access control list that indicates which members

of the enterprise are allowed to access and/or modify that particular resource, and the access control list information for each resource **115** is received along with each respective resource. Each resource **115** includes entities **210** that are information related to the enterprise. For example, a resource **115a1** of the enterprise that identifies the names and positions of those in management of the enterprise where “Jane Doe” may be in management of the enterprise in the position of “Vice-President of Sales,” may include “Jane Doe” as entity **210a** and “Vice-President of Sales” as entity **210b**.

[0031] Entity facts **220** of the entities **210** from the resources **115** of the enterprise are identified where each entity fact **220** describes at least one feature of the entity **210**, and the features include corresponding relationships between the entities **210** (**304**). For example, from the example above, an entity fact **220** may be a relationship between “Jane Doe” and “Vice-President of Sales,” where the entity fact **220** is “Jane Doe” is the “Vice-President of Sales.” As such, a feature of “Jane Doe” is “Vice-President of Sales,” and a feature of “Vice-President of Sales” is “Jane Doe.”

[0032] For each entity fact **220**, an entity fact access control list **230** is determined from the access control list of the resource from which the entities **210** of the entity fact **220** were derived (**306**). In the current implementation, the access control list of resource **115a1**, which included entities **210a** and **210b**, will be used as the entity fact access control list **230** of entity fact **220**. For example, if the access control list for resource **115a1** is the “Sales Department,” then the members of the enterprise that are included in the “Sales Department” group, as defined at the enterprise information system **110**, will have access privileges to the entity fact **220** of “Jane Doe” is the “Vice-President of Sales.”

[0033] Data describing the entities **210**, entity facts **220**, and the respective entity fact access control lists **230** are stored in data storage **134** (**308**). Indexer **136** then indexes the information in data storage **134**. The members of the enterprise may then use indexer **136** to, for example, search data storage **134** to access entity facts **220** of the resources **115** of the enterprise.

[0034] Each member of the enterprise is provided access privileges to the data describing the entities **210** and the entity facts **220** according to the respective entity fact access control list **230** in an indexer **136**. For example, in the example above, if “Alvin” were a member of the “Sales Department,” then Alvin would have access privileges to the entity fact **220** of “Jane Doe” is the “Vice-President of Sales” because Alvin is included in the entity fact access control list **230**.

[0035] Members of the enterprise can query the search engine **138** for information and/or one or more resources **115**. A query provided by a member of the enterprise will search indexer **136** for data, including resources **115**, entities **210**, and entity facts **220** that are stored in data storage **134**, which may be relevant to the member’s query. The entity fact access control lists **230**, in some implementations, are applied to the query provided by the member, and the member may only receive information, including resources **115**, entities **210**, and entity facts **220**, that the member has access privileges to according to the access control lists associated with the information. Additionally, in some implementations, member information may also be provided from the enterprise information system **110** to the data system

**130.** Member information may include, for example, the role or position of one or more of the members in the enterprise, an organizational structure of the enterprise, the contacts of one or more members in the enterprise (e.g., email or telephone), the member's relationship with the provided query, as described below.

**[0036]** Information may be determined to be relevant to the member's query based on how relevant a resource, entity, and/or entity fact is determined to be to a member's query. Additional analysis and criteria may be applied to the member's query to determine what is to be provided in response to the member's query. For example, the member's title, position, or group within the enterprise can factor into what is relevant to the member as well as the content, including the source information **112** and resources **115**, the member interacts with and uses. Additionally, activities of the member and the time and date the query was provided can factor into the relevancy. For example, interests related to the query member and information in the query member's email and calendar resources can factor into the relevancy. The search engine **138** may determine and analyze the relevancy to a member's query by the use of a knowledge graph; however, other methods of determining relevancy and relationships between resources **115**, entities **210**, and entity facts **220** may be used.

**[0037]** In the process of determining that an entity **210** or entity fact **220** is relevant to the query provided by a member, the search engine **138** also determines the query member's relationship with one or more entity **210** and entity fact **220**. In some implementations, the search engine **138** may first determine whether the provided query corresponds to, or otherwise is determined to be relevant to, at least one of an entity **210** and entity fact **220**. The search engine **138** then forms a relation between the provided query and the entities **210** and entity facts **220** determined to be relevant to the provided query. The search engine **138** may also determine the query member's relationship with the entities **210** and entity facts **220** determined to be relevant to the provided query.

**[0038]** FIG. 4 is a block diagram of a portion of an example enterprise knowledge graph representation **400** of enterprise relationship information derived by the integration system **132**. The knowledge graph has nodes and edges. Each node in the knowledge graph represents resources **115** and entities **210**, and pairs of nodes in the knowledge graph are connected by one or more edges. Each edge representing a relationship dimension that defines an entity fact **220** between two resources **115**, two entities **210**, or one of each of a resource **115** and an entity **210** represented by the pair of nodes, or several edges represent a series of relationships that connect two nodes by one or more intermediate nodes. As previously stated, the information to form an entity fact **220** between entities **210** is determined from the content provided in resources **115**. As shown in FIG. 4, the edges are unidirectional, but in other variations the edges may be bidirectional.

**[0039]** For example, the enterprise knowledge graph **400** includes nodes **410** and **414** that are connected by edge **412**. In the present example, node **410**, a member of the enterprise, "Jane Doe," is defined to be connected, by edge **412**, with node **414**, "Sales Department." Additionally, node **410** is determined to be connected, by edge **416**, with node **418**, "Project A." Node **422**, "Karen," another "Sales Department" employee, as shown by the edge **424** between node

**422** and node **414**, is also connected by edge **420** to node **418**. In the present example, node **410** is indirectly connected to node **422** via two separate relationships (node **414** and node **418**). Node **410** is also connected to another project, "Project B," at node **428** by edge **426**, and node **410** is connected to "Engineering Department," at node **432**, by edge **430**. For example, "Jane Doe," at node **410**, may have corresponded with the "Engineering Department" in the past to determine if a project was feasible or to check on the status of an event or project. As such, an information source **112** (e.g., enterprise email) would include a resource **115** (e.g., email message) that would include the entities of "Jane Doe" and "Engineering Department," where an entity fact **220** could describe a relationship between the entities. Further, another enterprise member, "Karen," who works in the Engineering Department, at node **436** is connected to node **432** by edge **434**.

**[0040]** Based on the exemplary enterprise knowledge graph **400**, if enterprise member "Jane Doe," the entity at node **410**, were to provide a search query "Karen," the Data System **130** would analyze the enterprise knowledge graph **400** to determine what is relevant to the query. For example, based on the connection between "Jane Doe" at node **410** and "Karen" at node **422** via node **418**, "Project A," and node **420**, "Sales Department," contact information or meeting options of "Karen" at node **422**, and resources **115** (e.g., shared documents or email messages) between identifications associated with "Jane Doe" and "Karen" may be determined to be relevant to the query provided. However, contact information or other information pertaining to "Karen" at node **436** may be determined to be less relevant or not relevant at all based on the connection shown between "Jane Doe" at node **410** and "Karen" at node **436**.

**[0041]** Additionally, based on the exemplary search query "Karen," the data system **130** may also determine information and resources **115** related to "Project A," shown at node **418**, are relevant to the query based on the connection of both node **410** and node **422** to "Project A." However, information and resources **115** related to "Project B," shown at node **428**, may not be determined to be relevant based on the lack of a connection between node **428** and node **422**.

**[0042]** Further, the enterprise knowledge graph **400** includes "Susan" at node **440** that is connected to "Jane Doe" at node **410** by edge **438**. "Susan" may be, for example, a sales lead that "Jane Doe" has previously met with and/or exchanged correspondence with. "Susan," in the current example, at node **440** is connected to "Company A" at node **444** by edge **442**, as she may be, for example, an account manager at "Company A." Additionally, node **444** is connected to "Clients" at node **448** of the enterprise by edge **446**, and node **448** is connected to the "Sales Department" at node **414** by edge **450**. For example, if "Mark," seen at node **458**, were another account manager for "Company A," and in Jane Doe's previous email exchange with Susan, she told Jane Doe that her sales proposal needed to be discussed with "Mark" for approval, edge **451** may be created between nodes **440** and **458** based on the entities provided in resources **115** (e.g., enterprise email).

**[0043]** Based on the exemplary enterprise knowledge graph **400**, if enterprise member "Jane Doe," the entity at node **410**, were to provide a search query "Company A," the Data System **130** would analyze the enterprise knowledge graph **400** to determine what is relevant to the query. For example, based on the connection between nodes **410** and

**440**, documents, emails, and biographical and contact information related to “Susan” may be determined to be relevant. Additionally, based on the exemplary sales lead email where Susan advised Jane Doe to speak to Mark, biographical and contact information for Mark may also be determined to be relevant.

[0044] Further, for example, Jane Doe has an upcoming trip scheduled to “Atlanta,” seen at node **454** and connected by edge **452** to node **410**. If, for example, the information related to Mark at node **458** indicates that Mark’s office is in “Atlanta,” then a connection can be made between Atlanta node **454** and Mark node **458** by edge **456**. As such, in the results provided to enterprise member Jane Doe’s query of “Company A,” the data system **130** may determine that it is relevant to provide enterprise member Jane Doe with an option to compose an email to Mark, call Mark, or schedule a calendar invitation to meet with Mark. As such, the entity facts **220** created based on entities **210** and resources **115**, in some implementations, can be opportunities for current or future action. The query results as described above are merely exemplary, and other types and methods of providing query results may be included. In some implementations, the enterprise knowledge graph **400** may include information related to an organizational structure of the enterprise, which can be obtained from one or more resource **115**, where the positions and members holding those positions are relationally structured in the enterprise knowledge graph **400**.

[0045] FIG. 5 is a flow chart of an example process **500** for providing search results to an enterprise member based on a query provided by the enterprise member. The process **500** can, for example, be implemented by the data system **130**. In some implementations, the operations of the example process **500** can be implemented as instructions stored on a non-transitory computer readable medium, where the instructions cause a data processing apparatus to perform operations of the example process **500**.

[0046] A search query provided by a member of the enterprise is received by the data system **130** (**502**). In some implementations, the enterprise information system **110** may provide an interface to the member of the enterprise to receive queries that are input by the member. The interface may be provided via, for example, an application or program that is in communication with the data system **130**, or in some implementations, the interface may communicate to the data system **130** directly.

[0047] Once the query is received, the search engine **138** may search the data storage **134**, which stores the data received from the enterprise information system **110** and that has been indexed (**504**). As previously described, the data system **130** receives resources **115**, including member information, and the resources **115** include data describing entities **210** and entity facts **220**, where the entity facts **220** are identified from the entities **210** of the resources **115**. Each entity fact **220** describes at least one feature of an entity **210**, where the features describe corresponding relationships between the entities **210**. Additionally, the indexer **136** includes data defining access privileges to the data entity fact access control list **230** for each entity fact **220**, which is determined from the access control list of each resource **115**, as previously described.

[0048] The entity facts **220** that are accessible to the member are determined based on the entity fact access control lists **230** (**506**), as previously described. For example, only facts and resources for having an access

control list that specifies the member has access to the fact or resource are provided. Other facts and resources to which the member does not have read access are not identified for the member.

[0049] Further, search result data including data describing entities **210** and entity facts **220** relevant to the query are determined based on member information of the member and entity facts **220** that are accessible to the member (**508**). For example, entities **210** and entity facts **220**, as previously described, which are determined to be relevant to the member’s query are identified. Additionally, the member’s position or relationships within the enterprise can factor into what is relevant to the member query as well as the source information **112** and resources **115** the member interacts with and uses. Activities of the member and the time and date the query was provided can determine or factor into the relevancy. Further, an enterprise knowledge graph representation, like enterprise knowledge graph **400**, may be used to determine the relevancy of resources **115**, entities **210**, and entity facts **220** of enterprise relationship information.

[0050] Search results to the query provided by the enterprise member are provided based on the search result data (**510**). The search result data may include entities **210**, entity facts **220**, the underlying resources **115** from which the entities **210** and/or entity facts **220** are derived, or a combination thereof. The underlying resources **115** from which the entities **210** and/or entity facts **220** are derived may be embedded in the entity **210** and/or entity fact **220** that is presented. For example, the underlying resource **115** may be provided as a link (e.g., hyperlink), drop down selection box, hover selection, among others. In some implementations, if an entity fact, for example **220a**, is derived from more than one resource, for example both resource **115a1** and resource **115b1**, then the embedded underlying resource for the entity fact **220a** may include one or both (or more) of the resources from which the entity fact **220a** was derived. In implementations where only one of the underlying resources **115** is provided, the selection of the underlying resource may be based on the most relevant resource of each of the underlying resources to the member. For example, if the query member is part of the sales department, and underlying resource **115a1** is stored in the sales department location of the database while underlying resource **115b1** is stored in the engineering department location of the database, resource **115a1** may be determined to be most relevant to the member. Additionally, in some implementations, query suggestions may be provided to the member based on the query that was provided or is in the process of being provided (e.g., the member is currently inputting) and/or the search results. The data system **130** may be configured to interpret and understand natural language input.

[0051] The presentation of the search results may be provided as a list of the search result data, and in some implementations one or more knowledge panel may be provided in the search results. A knowledge panel may be presented inline or adjacent to other search results for a received query or in place of the other search results. The knowledge panel may provide information related to an entity **210** and/or entity fact **220** referenced by a search query. In some implementations, a knowledge panel may provide a summary of information for the entity **210** and/or entity fact **220**. For example, a knowledge panel for “Jane Doe,” as provided above at node **410**, may include an enterprise profile picture of Jane Doe, the department that

she works in, her contact information, a biography, among other content (including a link to Jane Doe's enterprise webpage). A knowledge panel for a particular entity **210** and/or entity fact **220** may also, or alternatively, include information about other entities **210** and/or entity facts **220** that have been identified as related to the particular entity **210** and/or entity fact **220**. For example, if the enterprise member Jane Doe referenced above is the sales account manager for "Company A," the knowledge panel for the enterprise member Jane Doe may include information about Company A, like contact and location information. For example, the information about Company A may also include information about a sales lead contact at Company A, for example "Susan" (at node **440**). Information about Susan may also include a Company A profile picture of Susan, the department that she works in, her contact information, a biography, among other content (including a link to Susan's Company A enterprise webpage).

**[0052]** Entity Fact Search Processing

**[0053]** FIG. 6 is a flow chart of an example process **600** for ranking and providing the search results to an enterprise member based on a query provided by the enterprise member. The process **600** can, for example, be implemented by the data system **130**. In some implementations, the operations of the example process **600** can be implemented as instructions stored on a non-transitory computer readable medium, where the instructions cause a data processing apparatus to perform operations of the example process **500**.

**[0054]** In some implementations, the search result data may be ranked based on the relevance of the entities **210** and entity facts **220** in relation to the query and member information of the query member (**602**). As described above, the relevant entities **210** and entity facts **220** may be determined, and then results may be ranked or organized based on the level of relevance based on the query member and query provided. The search result data may then be provided in the order of the rankings to the member of the enterprise (**604**).

**[0055]** In some implementations, a quality score may be determined for each entity **210** and entity fact **220** based on the query provided by the member of the enterprise. The quality score may indicate the level of relevance for each entity **210** and entity fact **220** to the query. The quality score can be member specific, a partial or total aggregate, or a combination thereof. The enterprise knowledge graph, as described in FIG. 4, may be used and factored in to the quality score to associate the relevance of the entities **210** and entity facts **220** in relation to the query and member of the enterprise providing the query. Additionally, the quality score for the entities **210** and entity facts **220** can be based on explicit indications and/or implicit indications by the member of the enterprise or a collection of members of the enterprise. For example, a member of the enterprise may explicitly indicate they prefer information (e.g., entities **210** and entity facts **220**) that are from particular resources **115** or information sources **112** (e.g., the internet, email, or a sales department database). Other explicit indications may include a preference of resources **115** the member is an author of, resources **115** created by a member in the member of the enterprise's department or group, or interests expressed by the user (e.g., the stock market price of the enterprise), among others.

**[0056]** Implicit indications may be any type of indication that is provided without the member's explicit input. As already described above, in order to determine the relevancy

of entities **210** and entity facts **220**, different information related to the member's interactions with the enterprise system **110** and data system **130** can provide implicit indications. For example, implicit indications may include the member's position, group, or department within the enterprise, the information sources **112** and resources **115** the member has recently or regularly interacted with, and activities of the member.

**[0057]** Additionally, implicit indications can be provided by the member or a collection of members of the enterprise based on how the member or collection of members interacts with the search results that are provided to the query. For example, in some implementations, parameters may be applied to the quality score. Exemplary parameters may be one or more of a mean reciprocal rank, a mean precision, an average click position, a click through rate, and an abandonment rate, among others. The mean reciprocal rank can be a statistical measure for evaluating any process that produces a list of possible query results to a sample of queries, ordered by probability of correctness. The reciprocal rank of a query response is the multiplicative inverse of the rank of the first correct answer, and the mean reciprocal rank is the average of the reciprocal ranks of results for a sample of queries. Average click position may be the average rank in the order of provided query search results that the member selected. Click through rate may be the rate that a particular query search result is selected, and abandonment rate may be a number of abandoned queries divided by the number of queries related to one or more of the same or similar queries or the same or similar queries provided by a member a collection of members. An abandoned query is a query that has no subsequent query search result selection and does not have any subsequent refinements to the query. These are only exemplary parameters and other parameters or modifications to the parameters above may be provided.

**[0058]** Further, in some implementations, click result information can be included in the quality score. For example, if the member or average of collection of members, select a search result with a long click (e.g., the click lasting longer than a threshold period of time), then the search result in relation to the search query may be determined to be more relevant than if a search result is selected with a short click, where it could be determined that a certain percentage of time the search result selected with the short click was inadvertently selected. Additionally, search results that are not clicked to be selected may be determined to be less relevant to the query provided.

**[0059]** The relevancy, including the quality score, of each search result to each query provided may auto-tuned, that is modified and/or adjusted, by the data system **130** based on the member's or collection of member's interaction with the search results. For example, if the member provides a query for "Company A" and selects the second search result that is related to an entity fact **220** about "Susan" at Company A, then the quality score for that entity fact **220** may increase while quality scores for other entities **210** and entity facts **220** provided in the search results, that were determined to be relevant, may decrease or otherwise adjust. Additionally, enterprise acronyms and synonyms associated with entities **210** and entity facts **220** may be determined by the data system **130** based on the auto-tuning process.

**[0060]** In some implementations, the auto-tuning process may be performed automatically by the data system **130** without human input or intervention after the auto-tuning

process is established. Many enterprise data systems **110** contain sensitive and secret information that must remain confidential within the enterprise; therefore, intervention or input at the data system **130** may not be provided in order to preserve the confidentiality and secrecy of information within the enterprise. In some implementations, the relevancy determination, including the quality score, may be based on a static and/or dynamic algorithm, formula, heuristic, or a combination thereof.

**[0061]** Enterprise Assistants

**[0062]** The entity facts and related information can also be used for other information support operations for members of the enterprise. One example expanding the role of computer-implemented assistants to use enterprise information, subject to the access control lists.

**[0063]** FIG. 7 is an example environment **700** that incorporates computer-implemented assistants into the example environment structure of example environment **100**. Additionally, in some implementations, one or more members of the enterprise can use the computer-implemented assistants **710**, as seen in example environment **700**. As seen in FIG. 7, the one or more computer-implemented assistants **710a**, **710b**, . . . , **710N**, are communicatively connected to the enterprise information system **110**, data system **130**, member information system **720** (seen in FIG. 7 as **720a**, **720b**, . . . , **720N**), and with each other computer-implemented assistant **710**, or combinations thereof. However, such communication is not required, and in some embodiments, communication with more or fewer devices and systems may be included. Member information system **720** can include information that is not described or included in the resources **115**, but is accessible to the particular member associated with a particular computer-implemented assistant. For example, a member's personal email, personal calendar, or personal travel information may be included in the member information system **720**. The member information system **720** may be implemented on one or more devices (e.g., computer, mobile device, smart phone, tablet), and the storage of the member personal information may be on one or more of the devices, on different devices, or in a cloud storage environment, among others. The computer-implemented assistant **710** is configured to be implemented in any compatible software application, and the computer-implemented assistant **710** is not required to be used on only one application and/or one user device. Additionally, as seen in FIG. 7, in some implementations, the computer-implemented assistant **710** may communicate with data system **130** without the use of API **120**.

**[0064]** As previously described, data system **130** accesses and integrates enterprise resources **115** that include data describing entities **210** and entity facts **220**. Each entity fact **220** includes an entity fact access control list **230** that is provided from the underlying resource **115** from which the entity fact **220** was obtained, and each member has access privileges to resources **115**, entities **210**, and entity facts **220** according the respective entity fact access control lists **230**. Additionally, in some implementations, member personal information may also be provided from the member information system **720** to the enterprise information system **110** and/or the data system **130**.

**[0065]** The computer-implemented assistant **710** for each member of the enterprise can determine enterprise information specific to the member, including entities **210** and entity facts **220**, and integrate the member personal information

with the enterprise information specific to the member. For example, the computer-implemented assistant **710** can determine a calendar for the particular member based on the member personal information and the enterprise information specific to the member. Based on the computer-implemented assistant **710** combining this information, the computer-implemented assistant **710** can determine enterprise actions to perform for the member, where the enterprise actions are actions specific to the member's enterprise-related responsibilities. For example, the enterprise-related responsibilities include work meetings, arrangements, and travel, among others, but may not include scheduling a personal vacation.

**[0066]** In order to provide enterprise actions for the member, the computer-implemented assistant **710** may analyze and determine data related to one or more locations a user device has visited or is scheduled to visit, one or more enterprise member user devices that the user device has interacted with, interests of the member, industry information relevant to the member, and calendar information of the member, among others. This exemplary data may be provided by one or more user devices of the member, the enterprise information system **110**, the data system **130**, the member information system **720**, or a combination thereof.

**[0067]** Further, in some implementation, public resources **730** (seen in FIG. 7 as **730a**, **730b**, . . . , **730N**) that are published by entities external to the enterprise may be included in the information the computer-implemented assistant **710** uses to perform enterprise actions for the member. The computer-implemented assistant **710**, as previously described, may communicate with the data system **130**, and the schema provided at the data system **130** may also be used by the computer-implemented assistant **710** in order to sort and understand the information provided by the public resources. In some implementations, the API **120** is used to tag and index content and information received by the computer-implemented assistant **710**. Public resources may be, for example, information obtained on the internet, a travel database system, a professional networking application, among others. In the current implementation, the computer-implemented assistant **710** may determine a subset of public information from the public resources that is relevant to the particular member. For example, when the public resource is a travel database system and the member has a scheduled meeting in "Atlanta," the computer-implemented system **710** may determine that the subset of public information that is relevant to the particular member is the flights or other transportation modes from the current city the member is located to Atlanta, Ga. The computer-implemented assistant **710** may determine the subset of public information that is relevant to the particular member based on the member personal information, enterprise information specific to the member, and the entity facts **220**, among other information. After determining the subset of public information relevant to the particular member, the computer-implemented assistant **710** may determine enterprise actions to perform for the member.

**[0068]** Enterprise actions to perform for the member can be context specific and may include presenting insights or information to the member, presenting opportunities (e.g., a business lead) to the member, or a completion of one or more task for the member. For example, an enterprise action may include declining an invitation for a meeting. Declining could be based on a multitude of reasons, including the member having a vacation scheduled (even if this is only



listed on the member's personal calendar) or a scheduling conflict. Additionally, enterprise actions may include rescheduling a conflicting meeting, booking a flight, booking a hotel room, presenting local transportation options, finding restaurants suitable to the member, creating reservations at restaurants suitable to the member, presenting a daily agenda (where some items can be automatically generated), among others.

**[0069]** Additionally, an enterprise action may include providing relevant information (e.g., a document or slide presentation) to the member at the time of a scheduled meeting or event. The relevancy of information to a scheduled meeting or event may be based on the information available to the computer-implemented assistant **710**, including member personal information, enterprise information specific to the member, the entity facts, and the subset of public information. The computer-implemented assistant **710** may present to one or more of the user devices of the member the relevant information, including an enterprise resource, prior to or at the time of the event or meeting scheduled. For example, when the member has a scheduled meeting with "Tom," another member of the enterprise, and Tom has scheduled the meeting to discuss a particular document, the computer-implemented assistant **710** can determine the particular document is relevant to the meeting—for example, based on the title of the meeting, from an email exchange between Tom and the member, among others), and provide the document to the member at the time of the meeting or at a time prior to the scheduled meeting. Further, in some implementations, the computer-implemented assistant **710** may provide informational material, including public resources, to the one or more user devices of the member if that material is determined to be relevant to a scheduled meeting or event. For example, the member may have a meeting scheduled to discuss a complicated technical topic, and the computer-implemented assistant **710** may provide technical background information or other type of information to the member prior to the meeting.

**[0070]** In some implementations, the enterprise actions include presenting opportunities (e.g., a business lead) to the member. For example, if the member is in the sales department of the enterprise and has a scheduled trip to "Atlanta," where a sales lead is also visiting (e.g., known based on a previous email exchange between the member and the sales lead), the computer-implemented assistant **710** may present the one or more user devices of the member with a notification that the sales lead will also be in Atlanta at that time, and the computer-implemented assistant **710** can also ask or prompt the member if they would like to schedule a meeting or otherwise contact the sales lead while in Atlanta. After each enterprise action is determined to be performed, the computer-implemented assistant **710** causes the one or more user devices of the member to present a notification describing the enterprise action to the member.

**[0071]** The computer-implemented assistant **710**, in some embodiments, can coordinate between the computer-implemented assistant **710** of one or more other member of the enterprise. The computer-implemented assistant **710** of the member can determine that an enterprise action involves at least one other member of the enterprise. For example, if the computer-implemented assistant **710** of the member scheduled a business trip after planning a meeting with another member of the enterprise, the computer-implemented assistant **710** of the member can determine that the business trip

that has been scheduled affects the currently scheduled meeting. The computer-implemented assistant **710** may communicate with the computer-implemented assistant **710** of the other member and include information about the enterprise action that pertains to the at least one other member. For example, when the newly scheduled business trip conflicts with the scheduled business meeting, the computer-implemented assistant **710** of the member may provide information to computer-implemented assistant **710** of the other member that notifies the computer-implemented assistant **710** of the other member that the member is not able to make the meeting and other meeting times could be proposed. However, information from the enterprise action about where the member is traveling, who the member is meeting with, when the member is leaving for the trip, among other information related to the enterprise action, is not provided to the computer-implemented assistant **710** of the other member because that information would not pertain to the at least one other member. Further, in some implementations, the computer-implemented assistant **710** of each of the at least one or more other member may perform an enterprise action based on receiving the pertaining information of the enterprise action of the member.

**[0072]** Further, in some implementations, the computer-implemented assistant **710** is configured to be implemented on a third party application on the member device. For example, the third party application may be a travel system application that includes the member's boarding pass and other travel information. As such, the third party application can provide relevant third party application member personal information to the computer-implemented assistant **710**. Additional third party applications may also be used, for example, applications for banking, correspondence (e.g., email), professional networking, among others.

**[0073]** FIG. 8 is a flow chart of an example process **800** for providing enterprise actions by a computer-implemented assistant **710** for a member of the enterprise. The process **800** can, for example, be implemented by the computer-implemented assistant **710**. In some implementations, the operations of the example process **800** can be implemented as instructions stored on a non-transitory computer readable medium, where the instructions cause a data processing apparatus to perform operations of the example process **800**.

**[0074]** Enterprise resources **115** that are accessible to a member of the enterprise, based on the entity fact access control lists **230**, where the resources **115** include data describing entities **210** and relationships among the entities (including entity facts **220**), and enterprise information specific to the member are accessed (**802**). Additionally, as previously described, entity facts **220** of the entities **210** from the resources **115**, each entity fact **220** describing at least one feature of the entity **210** are identified (**804**). Further, as previously described, member personal information specific to the member, the member personal information including information not described by the enterprise resources that are accessible to the member is accessed (**806**). The computer-implemented assistant **710** then determines enterprise actions, as described above, to perform for the member based on the member personal information, enterprise information specific to the member and the entity facts (**808**). For each enterprise action determined to be performed, the computer-implemented assistant **710** provides a notification to the member describing the enterprise action performed (**810**).

**[0075]** Embodiments of the subject matter and the operations described in this specification can be implemented in digital electronic circuitry, or in computer software, firmware, or hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them. Embodiments of the subject matter described in this specification can be implemented as one or more computer programs, i.e., one or more modules of computer program instructions, encoded on computer storage medium for execution by, or to control the operation of, data processing apparatus. Alternatively or in addition, the program instructions can be encoded on an artificially-generated propagated signal, e.g., a machine-generated electrical, optical, or electromagnetic signal, that is generated to encode information for transmission to suitable receiver apparatus for execution by a data processing apparatus. A computer storage medium can be, or be included in, a computer-readable storage device, a computer-readable storage substrate, a random or serial access memory array or device, or a combination of one or more of them. Moreover, while a computer storage medium is not a propagated signal, a computer storage medium can be a source or destination of computer program instructions encoded in an artificially-generated propagated signal. The computer storage medium can also be, or be included in, one or more separate physical components or media (e.g., multiple CDs, disks, or other storage devices).

**[0076]** The operations described in this specification can be implemented as operations performed by a data processing apparatus on data stored on one or more computer-readable storage devices or received from other sources.

**[0077]** The term “data processing apparatus” encompasses all kinds of apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, a system on a chip, or multiple ones, or combinations, of the foregoing. The apparatus can include special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application-specific integrated circuit). The apparatus can also include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, a cross-platform runtime environment, a virtual machine, or a combination of one or more of them. The apparatus and execution environment can realize various different computing model infrastructures, such as web services, distributed computing and grid computing infrastructures.

**[0078]** A computer program (also known as a program, software, software application, script, or code) can be written in any form of programming language, including compiled or interpreted languages, declarative or procedural languages, and it can be deployed in any form, including as a stand-alone program or as a module, component, subroutine, object, or other unit suitable for use in a computing environment. A computer program may, but need not, correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub-programs, or portions of code). A computer program can be deployed to be executed on one computer or

on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

**[0079]** The processes and logic flows described in this specification can be performed by one or more programmable processors executing one or more computer programs to perform actions by operating on input data and generating output. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application-specific integrated circuit).

**[0080]** Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read-only memory or a random access memory or both. The essential elements of a computer are a processor for performing actions in accordance with instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto-optical disks, or optical disks. Devices suitable for storing computer program instructions and data include all forms of non-volatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

**[0081]** To provide for interaction with a user, embodiments of the subject matter described in this specification can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube) or LCD (liquid crystal display) monitor, for displaying information to the user and a keyboard and a pointing device, e.g., a mouse or a trackball, by which the user can provide input to the computer.

**[0082]** Embodiments of the subject matter described in this specification can be implemented in a computing system that includes a back-end component, e.g., as a data server, or that includes a middleware component, e.g., an application server, or that includes a front-end component, e.g., a user computer having a graphical user interface or a Web browser through which a user can interact with an implementation of the subject matter described in this specification, or any combination of one or more such back-end, middleware, or front-end components. The components of the system can be interconnected by any form or medium of digital data communication, e.g., a communication network. Examples of communication networks include a local area network (“LAN”) and a wide area network (“WAN”), an inter-network (e.g., the Internet), and peer-to-peer networks (e.g., ad hoc peer-to-peer networks).

**[0083]** While this specification contains many specific implementation details, these should not be construed as limitations on the scope of any inventions or of what may be claimed, but rather as descriptions of features specific to particular embodiments of particular inventions. Certain features that are described in this specification in the context

of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable subcombination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a subcombination or variation of a subcombination.

**[0084]** Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

**[0085]** Thus, particular embodiments of the subject matter have been described. Other embodiments are within the scope of the following claims. In some cases, the actions recited in the claims can be performed in a different order and still achieve desirable results. In addition, the processes depicted in the accompanying figures do not necessarily require the particular order shown, or sequential order, to achieve desirable results. In certain implementations, multitasking and parallel processing may be advantageous.

What is claimed is:

1. A method implemented in a data processing apparatus, comprising:

receiving resources of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members of the enterprise, and the resources including data describing entities related to the enterprise and relationships among the entities;

identifying entity facts of the entities from the resources of the enterprise, each entity fact describing at least one feature of the entity, wherein the features include corresponding relationships between the entities;

determining, for each entity fact, from the access control list of each resource, an entity fact access control list; storing data describing the entities, entity facts and the respective entity fact access control lists in a searchable index, wherein each entity fact is associated with its corresponding entity fact access control list; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists.

2. The method of claim 1, wherein each entity fact access control list is different from the access control lists provided for the resources of the enterprise.

3. The method of claim 2, wherein:

the members of the enterprise belong to a plurality of groups, each group including a respective subset of members of the members of the enterprise;

at least one entity fact access control list specifies access privileges for at least one group; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists comprises providing access to each member of the at least one group according to the access privileges.

4. The method of claim 1, further comprising:

receiving a query from a member of the enterprise;

receiving, in response to a search operation on the searchable index, search result data identifying entities and entity facts responsive to the query;

for each identified entity fact, determining whether the member has access to the entity fact based on the entity fact access control list;

for each identified entity fact for which the member has access to the entity fact based on the entity fact access control list, providing access to data describing the entity fact to the member; and

for each identified entity fact for which the member does not have access to the entity fact based on the entity fact access control list, not providing data describing the entity fact to the member.

5. The method of claim 4, wherein providing access to data describing the entity fact to the member comprises providing a knowledge panel display that displays the entity facts.

6. The method of claim 1, wherein determining, for each entity fact, from the access control list of each resource, an entity fact access control list comprises:

determining an entity fact of an entity from a resource; and

replicating the access privileges specified by the access control list for the resource in the entity access control list for the entity fact.

7. The method of claim 1, wherein at least one resource includes a plurality of access control lists, each of the plurality of access control lists corresponding to a subset of entity facts of the resource, and wherein each of the one or more members is provided access privileges to each of the subsets of entity facts for which the one or more members are included on the access control list corresponding to that respective subset of entity facts.

8. A system, comprising:

a processor; and

a computer-readable medium coupled to the processor and having instructions stored thereon, which, when executed by the processor, cause the processor to perform operations comprising:

receiving resources of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members of the enterprise, and the resources including data describing entities related to the enterprise and relationships among the entities;

identifying entity facts of the entities from the resources of the enterprise, each entity fact describing at least one feature of the entity, wherein the features include corresponding relationships between the entities;

determining, for each entity fact, from the access control list of each resource, an entity fact access control list;

storing data describing the entities, entity facts and the respective entity fact access control lists in a searchable index, wherein each entity fact is associated with its corresponding entity fact access control list; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists.

9. The system of claim 8, wherein each entity fact access control list is different from the access control lists provided for the resources of the enterprise.

10. The system of claim 9, wherein:

the members of the enterprise belong to a plurality of groups, each group including a respective subset of members of the members of the enterprise;

at least one entity fact access control list specifies access privileges for at least one group; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists comprises providing access to each member of the at least one group according to the access privileges.

11. The system of claim 8, further comprising:

receiving a query from a member of the enterprise;

receiving, in response to a search operation on the searchable index, search result data identifying entities and entity facts responsive to the query;

for each identified entity fact, determining whether the member has access to the entity fact based on the entity fact access control list;

for each identified entity fact for which the member has access to the entity fact based on the entity fact access control list, providing access to data describing the entity fact to the member; and

for each identified entity fact for which the member does not have access to the entity fact based on the entity fact access control list, not providing data describing the entity fact to the member.

12. The system of claim 11, wherein providing access to data describing the entity fact to the member comprises providing a knowledge panel display that displays the entity facts.

13. The system of claim 8, wherein determining, for each entity fact, from the access control list of each resource, an entity fact access control list comprises:

determining an entity fact of an entity from a resource; and

replicating the access privileges specified by the access control list for the resource in the entity access control list for the entity fact.

14. The system of claim 8, wherein at least one resource includes a plurality of access control lists, each of the plurality of access control lists corresponding to a subset of entity facts of the resource, and wherein each of the one or more members is provided access privileges to each of the subsets of entity facts for which the one or more members are included on the access control list corresponding to that respective subset of entity facts.

15. A computer-readable medium having instructions stored thereon, which, when executed by a processor, cause the processor to perform operations, comprising:

receiving resources of an enterprise, each resource having a respective access control list specifying access privileges to the resource for one or more members of the enterprise, and the resources including data describing entities related to the enterprise and relationships among the entities;

identifying entity facts of the entities from the resources of the enterprise, each entity fact describing at least one feature of the entity, wherein the features include corresponding relationships between the entities;

determining, for each entity fact, from the access control list of each resource, an entity fact access control list;

storing data describing the entities, entity facts and the respective entity fact access control lists in a searchable index, wherein each entity fact is associated with its corresponding entity fact access control list; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists.

16. The computer-readable medium of claim 15, wherein each entity fact access control list is different from the access control lists provided for the resources of the enterprise.

17. The computer-readable medium of claim 16, wherein:

the members of the enterprise belong to a plurality of groups, each group including a respective subset of members of the members of the enterprise;

at least one entity fact access control list specifies access privileges for at least one group; and

providing, to each of the members of the enterprise, access privileges to the data describing the entities and the entity facts in the searchable index according to the respective entity fact access control lists comprises providing access to each member of the at least one group according to the access privileges.

18. The computer-readable medium of claim 15, further comprising:

receiving a query from a member of the enterprise;

receiving, in response to a search operation on the searchable index, search result data identifying entities and entity facts responsive to the query;

for each identified entity fact, determining whether the member has access to the entity fact based on the entity fact access control list;

for each identified entity fact for which the member has access to the entity fact based on the entity fact access control list, providing access to data describing the entity fact to the member; and

for each identified entity fact for which the member does not have access to the entity fact based on the entity fact access control list, not providing data describing the entity fact to the member.

19. The computer-readable medium of claim 15, wherein determining, for each entity fact, from the access control list of each resource, an entity fact access control list comprises:

determining an entity fact of an entity from a resource; and

replicating the access privileges specified by the access control list for the resource in the entity access control list for the entity fact.

20. The computer-readable medium of claim 15, wherein the enterprise is a company, and wherein the enterprise information is information that is exclusively controlled by the enterprise.

\* \* \* \* \*