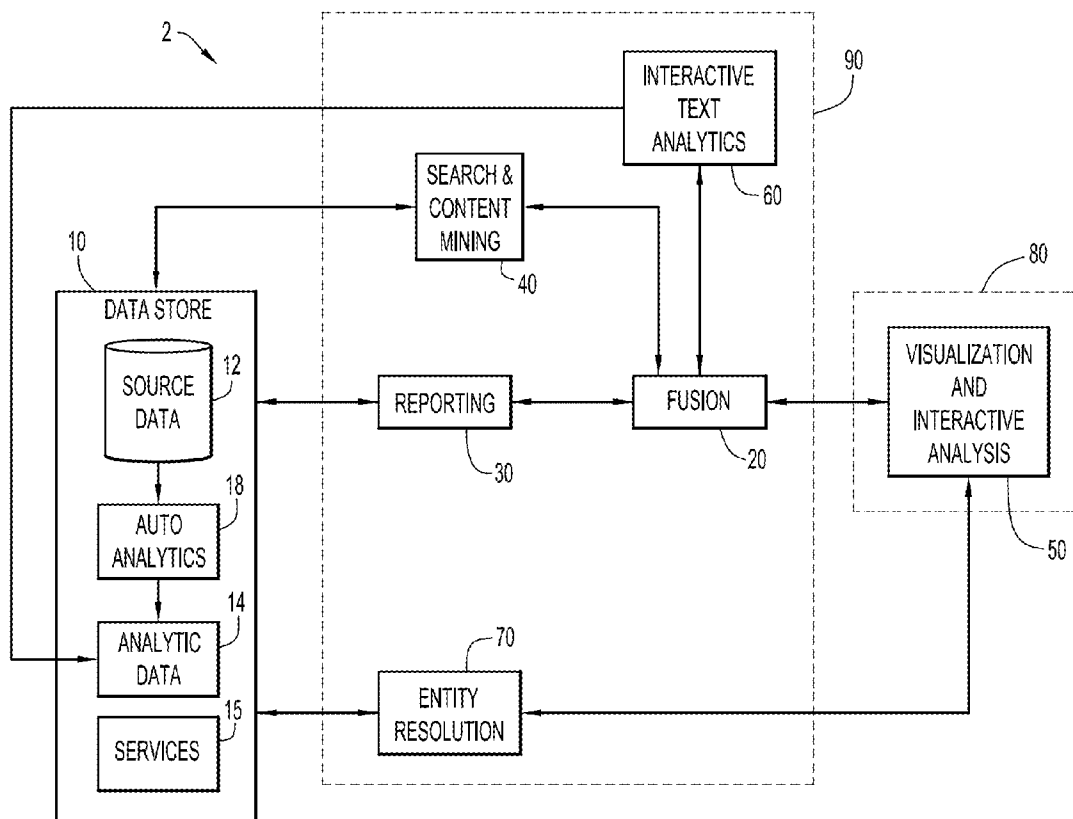




US 20160224629A1

(19) **United States**(12) **Patent Application Publication**
Beacom et al.(10) **Pub. No.: US 2016/0224629 A1**(43) **Pub. Date: Aug. 4, 2016**(54) **REDUCING A LARGE AMOUNT OF DATA TO
A SIZE AVAILABLE FOR INTERACTIVE
ANALYSIS****Publication Classification**(51) **Int. Cl.**
G06F 17/30 (2006.01)
(52) **U.S. Cl.**
CPC G06F 17/30486 (2013.01); **G06F 17/30554**
(2013.01)(71) Applicant: **International Business Machines
Corporation**, Armonk, NY (US)(72) Inventors: **Oliver J. Beacom**, Manotick (CA); **John
W. Bell**, Lees Summit, MO (US); **Paul
C. Jones**, Ottawa (CA); **James S. Luke**,
Isle of Wight (GB); **Luke Morgan**,
Cambridgeshire (GB); **Royston W.
Shufflebotham**, Cambridge (GB); **Mark
L. Todd**, Cambridge (GB)(21) Appl. No.: **14/983,885**(22) Filed: **Dec. 30, 2015****Related U.S. Application Data**(63) Continuation of application No. 14/609,699, filed on
Jan. 30, 2015.(57) **ABSTRACT**

A request for data is processed by identifying data within a data store satisfying the request, where an amount of the identified data exceeds a threshold, analyzing the identified data to produce an initial result set with a reduced amount of data including information summarizing the identified data, and generating one or more succeeding result sets each from a corresponding prior result set until a succeeding result set includes an amount of data satisfying the threshold. Each succeeding result set includes a reduced amount of data relative to the corresponding prior result set, and generating each succeeding result set includes filtering the corresponding prior result set and generating the succeeding result set including information summarizing the filtered result set.






FIG.1

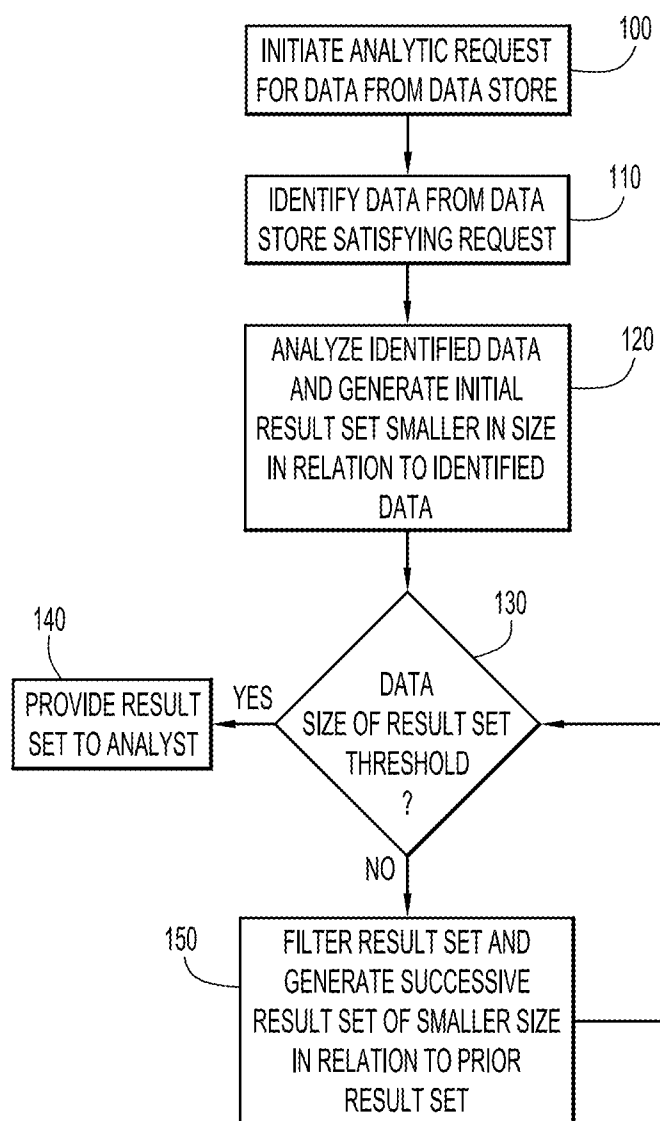


FIG.2





FIG.3






FIG. 5

200

212

205

208

210

211

213

214

215

216

217

218

219

220

221

222

223

224

225

226

227

228

229

230

231

232

233

234

235

236

237

238

239

240

241

242

243

244

245

246

247

248

249

250

251

252

253

254

255

256

257

258

259

260

261

262

263

264

265

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

295

296

297

298

299

300

301

302

303

304

305

306

307

308

309

310

311

312

313

314

315

316

317

318

319

320

321

322

323

324

325

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384

385

386

387

388

389

390

391

392

393

394

395

396

397

398

399

400

401

402

403

404

405

406

407

408

409

410

411

412

413

414

415

416

417

418

419

420

421

422

423

424

425

426

427

428

429

430

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

450

451

452

453

454

455

456

457

458

459

460

461

462

463

464

465

466

467

468

469

470

471

472

473

474

475

476

477

478

479

480

481

482

483

484

485

486

487

488

489

490

491

492

493

494

495

496

497

498

499

500

501

502

503

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

577

578

579

580

581

582

583

584

585

586

587

588

589

590

591

592

593

594

595

596

597

598

599

600

601

602

603

604

605

606

607

608

609

610

611

612

613

614

615

616

617

618

619

620

621

622

623

624

625

626

627

628

629

630

631

632

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647

648

649

650

651

652

653

654

655

656

657

658

659

660

661

662

663

664

665

666

667

668

669

670

671

672

673

674

675

676

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691

692

693

694

695

696

697

698

699

700

701

702

703

704

705

706

707

708

709

710

711

712

713

714

715

716

717

718

719

720

721

722

723

724

725

726

727

728

729

730

731

732

733

734

735

736

737

738

739

740

741

742

743

744

745

746

747

748

749

750

751

752

753

754

755

756

757

758

759

760

761

762

763

764

765

766

767

768

769

770

771

772

773

774

775

776

777

778

779

780

781

782

783

784

785

786

787

788

789

790

791

792

793

794

795

796

797

798

799

800

801

802

803

804

805

806

807

808

809

810

811

812

813

814

815

816

817

818

819

820

821

822

823

824

825

826

827

828

829

830

831

832

833

834

835

836

837

838

839

840

841

842

843

844

845

846

847

848

849

850

851

852

853

854

855

856

857

858

859

860

861

862

863

864

865

866

867

868

869

870

871

872

873

874

875

876

877

878

879

880

881

882

883

884

885

886

887

888

889

890

891

892

893

894

895

896

897

898

899

900

901

902

903

904

905

906

907

908

909

910

911

912

913

914

915

916

917

918

919

920

921

922

923

924

925

926

927

928

929

930

931

932

933

934

935

936

937

938

939

940

941

942

943

944

945

946

947

948

949

950

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

966

967

968

969

970

971

972

973

974

975

976

977

978

979

980

981

982

983

984

985

986

987

988

989

990

991

992

993

994

995

996

997

998

999

1000

1001

1002

1003

1004

1005

1006

1007

1008

1009

1010

1011

1012

1013

1014

1015

1016

1017

1018

1019

1020

1021

1022

1023

1024

1025

1026

1027

1028

1029

1030

1031

1032

1033

1034

1035

1036

1037

1038

1039

1040

1041

1042

1043

1044

1045

1046

1047

1048

1049

1050

1051

1052

1053

1054

1055

1056

1057

1058

1059

1060

1061

1062

1063

1064

1065

1066

1067

1068

1069

1070

1071

1072

1073

1074

1075

1076

1077

1078

1079

1080

1081

1082

1083

1084

1085

1086

1087

1088

1089

1090

1091

1092

1093

1094

1095

1096

1097

1098

1099

1100

1101

1102

1103

1104

1105

1106

1107

1108

1109

1110

1111

1112

1113

1114

1115

1116

1117

1118

1119

1120

1121

1122

1123

1124

1125

1126

1127

1128

1129

1130

1131

1132

1133

1134

1135

1136

1137

1138

1139

1140

1141

1142

1143

1144

1145

1146

1147

1148

1149

1150

1151

1152

1153

1154

1155

1156

1157

1158

1159

1160

1161

1162

1163

1164

1165

1166

1167

1168

1169

1170

1171

1172

1173

1174

1175

1176

1177

1178

1179

1180

1181

1182

1183

1184

1185

1186

1187

1188

1189

1190

1191

1192

1193

1194

1195

1196

1197

1198

1199

1200

1201

1202

1203

1204

1205

1206

1207

1208

1209

1210

1211

1212

1213

1214

1215

1216

1217

1218

1219

1220

1221

1222

1223

1224

1225

1226

1227

1228

1229

1230

1231

1232

1233

1234

1235

1236

1237

1238

1239

1240

1241

1242

1243

1244

1245

1246

1247

1248

1249

1250

1251

1252

1253

1254

1255

1256

1257

1258

1259

1260

1261

1262

1263

1264

1265

1266

1267

1268

1269

1270

1271

1272

1273

1274

1275

1276

1277

1278

1279

1280

1281

1282

1283

1284

1285

1286

1287

1288

1289

1290

1291

1292

1293

1294

1295

1296

1297

1298

1299

1300

1301

1302

1303

1304

1305

1306

1307

1308

1309

1310

1311

1312

1313

1314

1315

1316

1317

1318

1319

1320

1321

1322

1323

1324

1325

1326

1327

1328

1329

1330

1331

1332

1333

1334

1335

1336

1337

1338

1339

1340

1341

1342

1343

1344

1345

1346

1347

1348

1349

1350

1351

1352

1353

1354

1355

1356

1357

1358

1359

1360

1361

1362

1363

1364

1365

1366

1367

1368

1369

1370

1371

1372

1373

1374

1375

1376

1377

1378

1379

1380

1381

1382

1383

1384

1385

1386

1387

1388

1389

1390

1391

1392

1393

1394

1395

1396

1397

1398

1399

1400

1401

1402

1403

1404

1405

1406

1407

1408

1409

1410

1411

1412

1413

1414

1415

1416

1417

1418

1419

1420

1421

1422

1423

1424

1425

1426

1427

1428

1429

1430

1431

1432

1433

1434

1435

1436

1437

1438

1439

1440

1441

1442

1443

1444

1445

1446

1447

1448

1449

1450

1451

1452

1453

1454

1455

1456

1457

1458

1459

1460

1461

1462

1463

1464

1465

1466

1467

1468

1469

1470

1471

1472

1473

1474

1475

1476

1477

1478

1479

1480

1481

1482

1483

1484

1485

1486

1487

1488

1489

1490

1491

1492

1493

1494

1495

1496

1497

1498

1499

1500

1501

1502

1503

1504

1505

1506

1

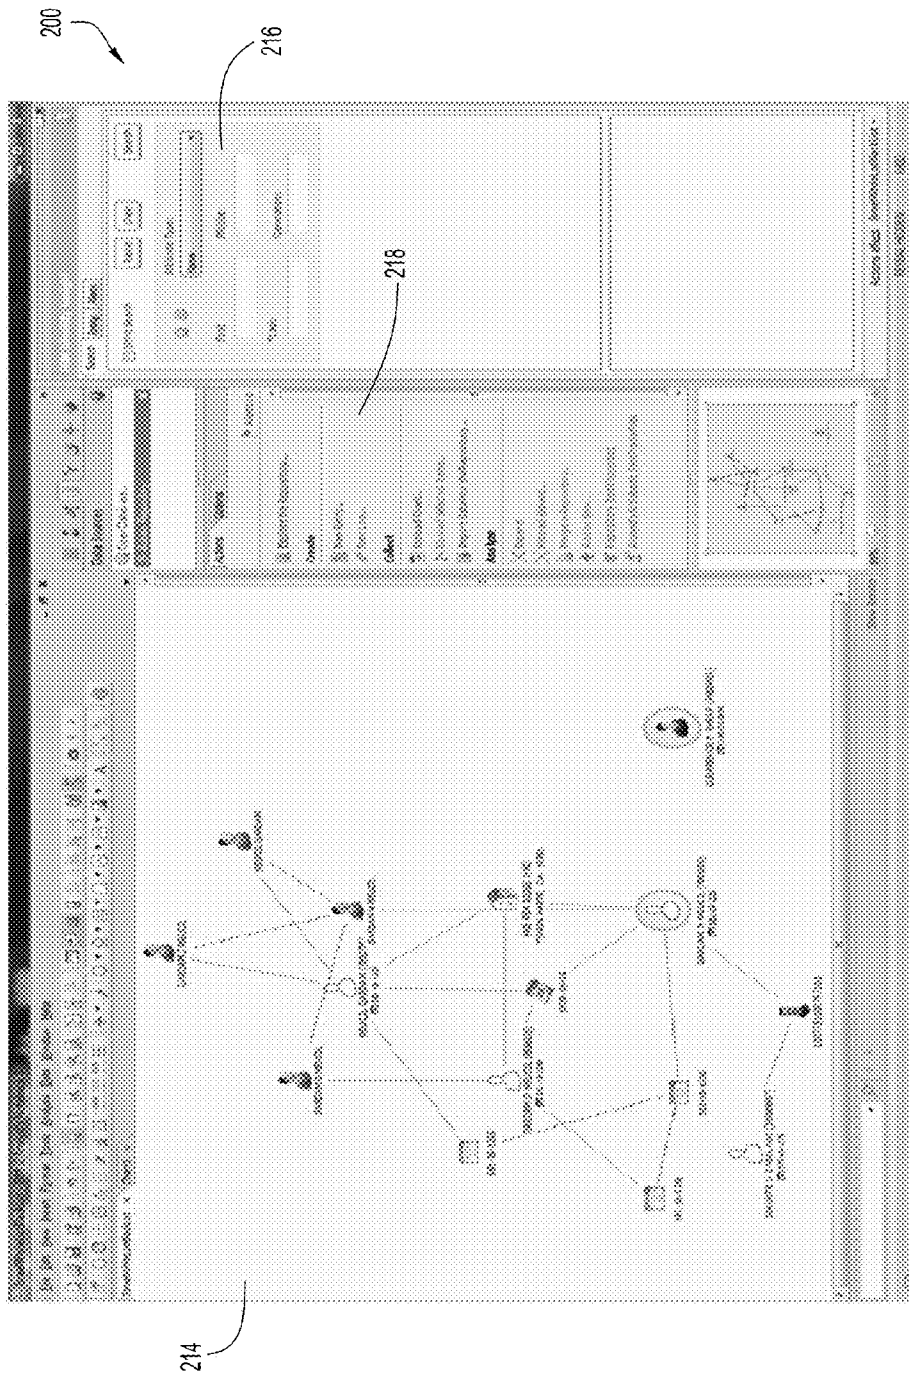


FIG.7

REDUCING A LARGE AMOUNT OF DATA TO A SIZE AVAILABLE FOR INTERACTIVE ANALYSIS

CROSS-REFERENCE TO RELATED APPLICATIONS

[0001] This application is a continuation of U.S. patent application Ser. No. 14/609,699, entitled “REDUCING A LARGE AMOUNT OF DATA TO A SIZE AVAILABLE FOR INTERACTIVE ANALYSIS” and filed Jan. 30, 2015, the disclosure of which is incorporated herein by reference in its entirety.

BACKGROUND

[0002] 1. Technical Field

[0003] Present invention embodiments relate to interactive analysis of a subset of data from a massive storage (e.g., peta-bytes) of data.

[0004] 2. Discussion of the Related Art

[0005] There are a number of industries that require analysis of large amounts of data, both structured and unstructured, in order to obtain meaningful access to a smaller subset of data that is of interest. For example, investigative agencies, such as law enforcement, intelligence and counter-fraud agencies, have access to very large (e.g., peta-bytes in data size) sources of data including call records, financial and/or computerized (electronic) transactions, etc. When combined with other conventional types of data, including unstructured data (e.g., intelligence reports), entity-link-property (ELP) data, it becomes a massive task for an analyst to handle such massive amounts of data to obtain a meaningful subset of the data for a particular search or analysis. In particular, consider typical sources of call data records, which may include a number of years of data for several million individuals, resulting in excess of one trillion items of data.

[0006] Such data cannot be processed for visual, interactive analysis due to the size of result sets obtained from the data source. For example, result sets that need to be analyzed, while smaller than the overall data source, are still too large to be stored in the memory of a conventional desk top or other computer or for analysis by conventional data analysis tools. In addition, the result sets that need to be analyzed cannot be visualized using conventional techniques, and the result sets are also too large to be transferred between computing devices or nodes due to a lack of available network bandwidth that would be required for such transfers.

SUMMARY

[0007] According to one embodiment of the present invention, a computer-implemented method of processing a request for data comprises identifying data within a data store satisfying the request, wherein an amount of the identified data exceeds a threshold, analyzing the identified data to produce an initial result set with a reduced amount of data including information summarizing the identified data, and generating one or more succeeding result sets each from a corresponding prior result set until a succeeding result set includes an amount of data satisfying the threshold, where each succeeding result set includes a reduced amount of data relative to the corresponding prior result set. The generation of each succeeding result set includes filtering the corresponding prior result set and generating that succeeding result set including information summarizing the filtered result set.

[0008] Embodiments of the present invention further include a system and computer program product for processing a request for data including features similar to those described above.

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

[0009] Generally, like reference numerals in the various figures are utilized to designate like components.

[0010] FIG. 1 is a diagrammatic illustration of an example system architecture for implementing an embodiment of the present invention.

[0011] FIG. 2 is a flow diagram illustrating processing of a large amount of data for interactive analysis according to an embodiment of the present invention.

[0012] FIGS. 3-7 provide graphical illustrations of a user interface that provides visualization of data in order to facilitate a reduction in the size of the data for further analysis according to an embodiment of the present invention.

DETAILED DESCRIPTION

[0013] Embodiments of the present invention enable real-time interaction with large amounts (e.g., peta-bytes or larger) of data to allow an analyst to filter result sets that can be analyzed by visual analytic tools. In particular, embodiments of the present invention facilitate processing of a request by an analyst for an analytic (e.g., a request for data based upon a search query, a find path request for finding a shortest path of all possible paths between two or more nodes of structured data, a request for geo-spatially related data (e.g., find entities within a defined area having a similar attribute such as the same type of credit card, etc.) within a data store comprising large amounts of data. A subset of data is identified within the data store that satisfies the request. However, due to the large size of data within the data store, the data subset can be greater than a threshold for allowing the data to be analyzed with data analysis tools available to the analyst (e.g., at a computing device of the analyst). Identified information within the data subset is summarized according to one or more facets to produce an initial result set with a reduced amount of data including information summarizing the identified data based upon the one or more facets. The analyst, utilizing the available data analysis tools, can filter the identified data so as to generate one or more succeeding result sets (where each succeeding result set is generated from a corresponding prior result set, and each succeeding result set includes information summarizing a reduced amount of identified data relative to the corresponding prior result set) until a succeeding result set is generated that includes information summarizing identified data to an amount that is at or below the threshold. Each succeeding result set that is generated based upon a filtering of the corresponding prior result set includes information summarizing the succeeding result set.

[0014] Embodiments of the present invention facilitate identified information to be provided for use by one or more analysis tools of the analyst only when the size of the filtered identified information is reduced to a level that is at or below the threshold. However, the embodiments of the present invention facilitate an analysis of a representation of the identified information based upon the facets associated with the identified information. This enables the analyst to visualize and further reduce the identified information, using the visu-

alization tools in accordance with the present invention, so as to obtain a meaningful result set that has been sufficiently filtered to a data size that permits a more detailed analysis of the data.

[0015] Embodiments of the present invention are useful for a number of applications in which large amounts of data need to be analyzed for a particular scenario. For example, consider law enforcement or other investigative agencies that need to peruse call data records within densely populated areas over a designated time period in order to find information pertaining to a particular investigation. The number of call records that might be within a data store available for searching, depending upon the population density for a particular area (e.g., New York City, Los Angeles, Chicago, etc.) and a selected time period (e.g., days or months), could easily number in the trillions. The embodiments of the present invention facilitate visualization by a data analyst of a reduced size of the initially identified search result information, where the size reduced information is visualized within a user interface for the analyst based upon facets associated with the search result information, and the analyst can filter the search result information (e.g., in an iterative process, with one or further filter operations being performed) until the search result information has been sufficiently reduced to a manageable size to be analyzed in greater detail by the analyst (utilizing analysis tools).

[0016] An example system architecture for implementing embodiments of the present invention is illustrated in FIG. 1. In particular, the system 2 includes a scalable data store 10 comprising source data 12 and analytic data 14. The data store 10 can include any suitable number of databases comprising any one or more suitable computing devices providing data storage and one or more suitable processors configured to execute instructions associated with one or more software modules that facilitate performance of database operations including parallel database operations (such as extract, transform and load operations) as well as other analytic processing operations of the stored data.

[0017] The data stored by the data store 10 includes both structured and unstructured data, where structured data can be obtained by analysis of the unstructured data. For example, the data store 10 can include structured data such as transactional data (e.g., call data records, financial transactions, computer/internet transactions, etc.), documents data (e.g., documents such as reports, graphs, etc. containing unstructured data of interest), entity data (e.g., data relating to persons or individuals) and entity-link-property (ELP) data. Unstructured data can be analyzed as described herein to extract entity and relationship data from the unstructured data so as to define facets of the data, where the facets are then used to summarize the identified data from an analytic request to generate result sets including data summary reports. The summary reports can be provided in a visualized format at an interactive user interface of the computing device of an analyst to facilitate filtering of result sets so as to reduce the amount of identified information to a manageable level.

[0018] The data store 10 further includes a data store services module 15 including any suitable number of software applications that facilitate processing of the data (e.g., performing analytic services to identify and associate data, finding shortest paths or find path techniques within graph analytics associated with the data, finding common neighbors within graph analytics associated with the data) as well as

other services associated with the data (e.g., services supporting searching of data within the data store 10).

[0019] An auto analytics module 18 provided within the data store 10 includes one or more software programs that analyze the source data 12 within the store (e.g., by running automated batch jobs on the data) and generate analytic data 14 within the store based upon algorithms for identifying relationships between elements or instances of the source data. For example, the auto analytics module 18 can automatically generate summary graphs of source data, where the summary graphs can then be searched utilizing the services module 15 (e.g., using find path techniques for the summary graphs). An example of a summary graph or source data is in relation to a collection of call records, where the summary graph includes data nodes and links between nodes based upon relationships between the nodes. The nodes can represent one or more calls, including a single node representing multiple calls (e.g., multiple calls between the same people, such as caller A and caller B). The analytic data 14 can be used to generate facets for summarizing the identified information of an analytic request within a result set.

[0020] In an example embodiment, the data store 10, including modules of the data store, is implemented utilizing a software platform commercially available from IBM Corporation under the trademark Pure Data for Operational Analytics (PDOA).

[0021] The system 2 also includes any suitable number of computing devices that communicate and exchange data with the data store 10 via any suitable one or more wired and/or wireless networks. The network(s) may be implemented by any number of any suitable communications media (e.g., wide area network (WAN), local area network (LAN), Internet, Intranet, etc.). Each of the computing devices includes one or more processors configured to perform operations as described herein based upon instructions provided by software applications within software modules and/or platforms stored within memories of the computing devices.

[0022] The system 2 includes a fusion module 20 with one or more suitable software applications that facilitate analyst operations of searching of the data store 10 and providing search results to a computing device of the analyst, via a visualization and interactive analysis module 50 as described herein. All search and/or other analytic requests and search results as described herein are routed for processing by the fusion module 20. The fusion module 20 facilitates collection and consolidation or fusion of data result sets for presentation to the analyst via module 50 in a useful manner. In an example embodiment, the fusion module 20 is implemented utilizing software commercially available from IBM Corporation under the trademark Intelligence Analysis Platform (TAP).

[0023] The fusion module 20 utilizes a reporting module 30 to execute search and/or other analytic operations for obtaining data from data store 10 (e.g., via interactions between the reporting module 30 and the services module 15 of the data store 10). As described herein, the reporting module 30 returns data from a search or other analytic request as one or more summary reports to the fusion module 20. The summary reports provided by the reporting module 30 to the fusion module 20 enable an analyst to refine the search results and/or analytic parameters of the search results in order to efficiently reduce or filter the size of the identified information originally obtained from the analytic request to a manageable size for further analysis and processing by the analyst (e.g., at the computing device of the analyst). The reporting module 30

summarizes data identified by a search or other analytic request utilizing facets associated with the identified data in order to generate the summary reports representing a visualization of the data at a smaller data size. The facets comprise information about relationships between instances of data, which can be obtained using the analytic data **14** generated within the data store by the auto analytics module **18** as well as other modules as described herein. In an example embodiment, the fusion module **20** is implemented based upon software commercially available from IBM Corporation under the trademark Cognos.

[0024] A search and content mining module **40** interacts with the fusion module **20** and data store **10** and includes one or more suitable software applications that facilitate searching of the data store **10** and also unstructured content mining, entity and relationship extraction against unstructured content for data within the data store **10**. In an example embodiment, module **40** crawls unstructured content of the source data **12** and applies text analytics annotators according to any one or more suitable algorithms to extract entity and relationship data into a structured form. The structured form of this data can be used to generate analytic data **14** within the data store **10** and/or used by other modules within the system **2** (e.g., the reporting module **30** and/or the entity resolution module **70**). The search and content mining module **40** can also enable analysts (operating, e.g., via visualization and interactive analysis module **50**) to input information that identifies correlations and links instances of data within the source data **12**, such as an identification of an association between entity data (e.g., Abraham Lincoln) and document data (e.g., "Gettysburg Address"). Module **40** enables a selection of a subset of unstructured content, such as unstructured documents, and further the limiting of analysis of data extracted from unstructured content to that which originates from the selected subset (e.g., analysis of extracted entity and relationships data from the unstructured content is limited to what has originated from the selected subset). In an example embodiment, the search and content mining module **40** utilizes a software platform commercially available from IBM Corporation under the trademark Watson Content Analytics to perform such operations.

[0025] The system **2** includes a visualization and interactive analysis module **50** that interacts with the fusion module **20** and also facilitates interaction of a user or analyst within the system. In particular, module **50** includes one or more suitable software applications that provide a user interface for the analyst (e.g., a user interface as depicted in FIGS. 3-7). The visualization and interactive analysis module **50** enables an analyst to initiate searches and other analytic requests by the fusion module **20** as well as filter the search requests based upon visualized displays of generated search reports as described herein. In an example embodiment, the visualization and interactive analysis module **50** is implemented utilizing software commercially available from IBM Corporation under the trademark IBM® i2® Analyst's Notebook®.

[0026] An interactive text analytics module **60** includes one or more suitable software applications that interact with the fusion module **20** and the data store **10** to facilitate manual analysis of source data **12** so as to generate analytic data **14** (e.g., by identifying relationships between one or more types of unstructured data). For example, an analyst using the visualization and interactive analysis module **50** and fusion module **20** can analyze source data and manually generate and/or validate analytic data that provides associations or links

between related data (e.g., associating entity data for two entities, such as person A and person B, based upon a known relationship of which the analyst is aware).

[0027] The system further includes an entity resolution module **70** that links the visualization and interactive analysis module **50** directly with the data store **10**. This module **70** includes one or more suitable software applications that read entity data from the data store **10** and perform entity and relationship resolution operations. For example, the entity and relationship resolution operations can facilitate the generation of entity graphs (and/or the resolution of already formed entity graphs) that link entity data with other types of data within the data store, where the entity graphs can be searched by an analyst via the visualization and interactive analysis module **50** and/or utilized by data store **10** (e.g., the data store services module **15**) for providing further analytic operations with the data stored therein. In an example embodiment, the entity resolution module **70** analyzes two or more data elements or instances of data (e.g., two or more data records, two or more rows of data within a table, etc.) to determine whether the instances of data represent the same entity. When two or more instances of data are determined as representing the same entity, these instances of data can be consolidated as a single entity for the summary reports generated by the reporting module **30**. This process can be performed manually by an analyst (via the visualization and interactive analysis module **50**) and/or automatically by the entity resolution module **70**.

[0028] The entity resolution module **70** is further configured to export entity graph data and/or other types of data back to the data store **10**. This facilitates access and use of resolved entity graphs by the data store **10** as well as other system modules (e.g., utilizing analytic tools such as find path tools) in addition to updating the analytic data within the data store (by resolving entity issues such as reducing multiple data elements for the same entity). A resolved entity graph generated by module **70** can significantly reduce clutter in a graph, e.g., by collapsing multiple instances of the same entity (e.g., the same person using different aliases) into a single entity within the graph. Module **70** can be configured to provide automated updates in resolved entity graphs for the data store **10** or, alternatively, provide such updates which are manually reviewed by an analyst prior to implementing resolved entity graphs into the data store **10**.

[0029] In an example embodiment, the entity resolution module **70** is implemented utilizing software commercially available from IBM Corporation under the trademark Infosphere Identity Insight.

[0030] Thus, the system **2** facilitates analysis of the source data **12**, both automatically (via the auto analytics module **18** and the entity resolution module **70**) as well as manually by an analyst (utilizing the visualization and interactive analysis module **50** in combination with the interactive text analytics module **60** and/or the entity resolution module **70**), to generate analytic data **14** so as to generate structured data by extracting entities and relationships from the unstructured data which facilitates adequate searching of the data store **10** to produce useful search results.

[0031] The analytic data **14** that is obtained further facilitates the summarization of data by the reporting module **30** so as to present summary reports in a chart or graph-based format that can be visualized using the user interface implemented by the visualization and interactive analysis module **50**. In particular, the reporting module **30** summarizes iden-

tified data within the data store **10** from a search or other analytic request based upon facets associated with the identified data, where the facets define relationships between instances of data based upon information such as the analytic data **14** within the data store **10**. As an example, consider identified data based upon a search request that includes call data records for a city over a specified period of time. An initial summary report generated by the reporting module **30** for the call records may comprise a graph with points or nodes representing entities with links between nodes representing calls between entities. Facets utilized to generate the summary report can include, e.g., information obtained from the analytic data **14** providing to relationships between entities, where multiple calls between two or more entities may be summarized as a facet comprising a single link between the entities. Other facets can also be utilized to define other types of relationships between entities and/or other types of data for purposes of generating a summary report of identified data by the reporting module **30**.

[0032] The system modules can be stored and implemented within computing devices and/or other hardware components in any suitable manner within the system **2**. As previously noted, the auto analytics module **18** can be implemented within computing devices and/or other hardware component(s) of the data store **10**. Any one or more of modules **30**, **40**, **50**, **60** and **70** can be implemented along with the fusion module **20** within one or more of the same computing devices or in any other one or more other computing devices, where each one or more computing devices comprise one or more memory structures to store module(s) as well as one or more processors configured to perform operations in accordance with computer readable program instructions provided by the software application(s) of the modules. In an example embodiment, the visualization and interactive analysis module **50** can be implemented within one or more computing devices of the analyst (e.g., a personal computer of the analyst, represented as dashed line element **80** in FIG. 1), which is separate from one or more computing devices implementing the fusion module **20**, reporting module **30**, search and content mining module **40**, interactive text analytics module **50** and entity resolution module **60** (represented as dashed line element **90** in FIG. 1). As previously noted herein, each computing device can communicate with other computing devices via any suitable wired or wireless network(s) (i.e., as indicated by the connecting arrows between modules **20-70** and the data store **10** as shown in FIG. 1).

[0033] Operation of the system **2** to provide search results to an analyst is now described with reference to the flow chart of FIG. 2. At **100**, a search or analytic request for data from the data store **10** is initiated by an analyst using the visualization and interactive analysis module **50** (e.g., implemented as a user interface at a computing device of the analyst). The visualization and interactive analysis module **50** communicates the request to the fusion module **20** for processing of the request. As previously noted, the data store **10** includes source data **12** and analytic data **14** (generated via the auto analytics module **18**, search and content mining module **40**, interactive text analytics module **60**, and/or entity resolution module **70**), and both types of data are searched at **110** (utilizing the search and content mining module **40**) to identify data within the data store **10** that matches the search or analytic request.

[0034] At **120**, the identified data is analyzed, utilizing the reporting module **30**, and an initial, result set is formed or generated utilizing facets associated with the identified data.

The initial result set summarizes and represents the identified data but is smaller in data size. For example, if the size of the identified data is on the order of gigabytes or greater (e.g., terabytes or petabytes), a result set that summarizes the identified data can be generated that is on the order of megabytes or smaller. The result set is generated that can be visualized to represent the identified data in a useful manner. For example, the result set can include one or more data summary reports that facilitate a suitable visualization of the identified data, such as a representation of the identified data in the form of a chart or a graph. Since the size of the identified data (particularly for data stores that store massive amounts of data) may be so large that the data cannot be analyzed utilizing the analysis tools of the analyst, the identified data is not initially provided to the fusion module **20**. Instead, the one or more summary reports of the result set are provided to the fusion module **20**.

[0035] At **130**, the fusion module **20** determines whether the data size of the result set representing or summarizing the identified data is no greater than (i.e., less than or equal to) a threshold value (e.g., a data size value that is a maximum size that can be processed by the fusion module **20** and/or the visualization and interactive analysis module **50**). If the data size of the result set is no greater than the threshold value, the result set (including the identified data represented by the result set) is provided at **140** to the fusion module **20** and/or the visualization and interactive analysis module **50** for further processing.

[0036] Alternatively, if the data size of the result set is greater than the threshold value, the result set is filtered at **150** to generate a successive result set of smaller size in relation to the original and/or prior result set. Filtering of the result set can be accomplished, e.g., via user input by the analyst via the visualization and interactive analysis module **50**. For example, one or more data summary reports generated by the reporting module **30** can be visualized as graphical displays of data rendered within a user interface by the visualization and interactive analysis module **50** (such as the user interface displays depicted in FIGS. 3-7). The analyst can selectively filter the result set based upon a selection of a number of different filter factors, such as a date range, a geographic area or region, specific types of records (e.g., cellular call records), etc. associated with the identified information of the result set. The filter factors are provided to the reporting module **30**, and the reporting module generates a further or successive result set that is smaller in data size in relation to the previous result set based upon the filter factors selected by the analyst.

[0037] As indicated by FIG. 2, the process steps of **130** and **150** are repeated in an iterative manner until the data associated with the successive result set that summarizes the identified information of the analytic request is no greater than the threshold value. At such point where the result set represents a data size that is no greater than the threshold value, identified data associated with the information of the result set is provided for analysis by the fusion module **20** and/or the visualization and interactive analysis module **30**.

[0038] An example embodiment utilizing the system and methods of the present invention is now described with reference to the user interface displays depicted in FIGS. 3-7. In this example embodiment, call data records are analyzed within a large metropolitan city. The user interface **200** depicted in FIGS. 3-7 is generated by the visualization and interactive analysis module **50** at an analyst's computing device based upon data summary reports generated by the

reporting module **30** and provided to the fusion module **20**. In this example embodiment, the analyst requests a search of the data store **10** for call records and related information, including cell tower locations, for New York City. The data store **10** includes about **1** trillion call data records that are accessible for search and analysis. An example of how information from the data summary reports for the initial result set are visualized is depicted within the user interface **200** of FIG. **3**. The user interface **200** comprises an interactive display including a graphical representation provided within a panel or window **202** of the user interface **200**, where the graphical representation provides a mapped view of the city (e.g., around a train station or other densely populated area). The analyst may be interested, e.g., in determining which entities made cell phone calls during a particular time period in the depicted area of the city. An analytic (e.g., search) request is initiated by the analyst (via the visualization and interactive analysis module **50**) that provides search terms in relation to this request. Another window **204** of the user interface **200** provides fields for input of information by the analyst, including information that may be used to filter identified information uncovered from the data store **10** in relation to the analytic request.

[0039] A portion of the mapped view within window **202**, such as region **203**, can be selected by the analyst utilizing the user interface tools of module **50**. Selection of this portion by the analyst comprises a filtering operation which generates a result set having a data size that is smaller than the result set associated with the visualized summary reports shown in FIG. **3**.

[0040] Referring to FIG. **4**, the filtered result set shows the number of call records for the selected region **203** over a year, where the calls per month are represented by bars **206** in a chart. The data field **205** in the user interface **200** allows the analyst to select different views of the data (e.g., viewing call volume in FIGS. **4** and **5** vs. viewing a grid of specific calls in FIG. **6**). A number of different filtering options are provided in the user interface **200** depicted in FIG. **4**, including a selection to display calls in the region **203** on a daily basis, weekly basis, etc. as well as types of calls/communications and other features in window **208**. The filtered results indicate that over 148 million calls (of the trillion+calls stored within the data store **10**) are in the identified information from the data store in the filtered result set. The data size of this result set may still be too large for analysis via the visualization and interactive analysis module **50** (i.e., the size of the result set is greater than a specified threshold value), such that further filtering of the data is required. However, the summary reports visualized via the user interface **200** (which is implemented with module **50**) are of a smaller size than the data size corresponding with the current result set, which enables a relatively fast and efficient mechanism to allow the analyst to further filter the data within the result set.

[0041] Selection by the analyst (using the features of module **50**) of a single month (e.g., November) from the result set visualized in FIG. **4** results in a further filtering of the result set to obtain the successive result set visualized by the user interface **200** as depicted in FIG. **5**. In this current result set, the number of calls has been filtered down to around 12 million.

[0042] The analyst can continue to filter the result set (e.g., if the current result set includes data that is still greater than the threshold value). For example, the analyst can select a specific day to view the specific call records associated with

the region **203** on the particular day. Visualization of the filtered result set is depicted in the user interface **200** of FIG. **6**, which indicates the result set includes 735 calls and with data for specific calls being visualized in a grid format **212**. Thus, the call records data of about 1 trillion calls in the data store **10** is searched and filtered down to a manageable and meaningful number (735) with relative ease.

[0043] Referring to FIG. **7**, the analyst can also have the current result set (i.e., 735 call records as depicted in FIG. **6**) visualized in different formats. For example, some or all of the call records for the selected day in the selected region **203** of the city can be visualized in a node map or graph linking nodes each other to represent calls made between the entities (e.g., individuals making the calls) or point locations (e.g., a house or dwelling). Searching of data within and/or further filtering of the current result set is enabled via input by the analyst within field of window **216**, while window **218** enables the analyst to search other data sources and/or perform other data analysis operations. In this example, the current result set includes data that is of sufficiently reduced size (the data size is no greater than the threshold value) such that the entire data for the current result set can now be provided from the data store **10** to the one or computing devices associated with the analyst for further analytic operations to be performed by the fusion module **20** and/or the visualization and interactive analysis module **50**.

[0044] Thus, embodiments of the present invention facilitate real-time interaction by an analyst (e.g., utilizing software tools at the computing device of the analyst) of massive data results (e.g., peta-bytes or larger) obtained from a data store as a result of an analytic request by the analyst. Since summary reports of each result set are utilized to visualize data within a user interface operated by the analyst, filtering of each successive result set can be performed rapidly and efficiently until the data size of a successive result set is no greater than a specified threshold value. The entire data of a result set is not provided to the computing device(s) operated by the analyst until the data size is no greater than the threshold value.

[0045] It will be appreciated that the embodiments described above and illustrated in the drawings represent only a few of the many ways of implementing embodiments for reducing a large amount of data to a size available for interactive analysis.

[0046] The environment of the present invention embodiments may include any number of computer or other processing systems (e.g., client or end-user systems, server systems, source systems, target systems, etc.) and databases or other repositories arranged in any desired fashion, where the present invention embodiments may be applied to any desired type of computing environment (e.g., cloud computing, client-server, network computing, mainframe, stand-alone systems, etc.). The computer or other processing systems employed by the present invention embodiments may be implemented by any number of any personal or other type of computer or processing system (e.g., desktop, laptop, PDA, mobile devices, etc.), and may include any commercially available operating system and any combination of commercially available and custom software (e.g., browser software, communications software, server software, database management system, analysis modules, data quality profiler module, data quality engine, data quality reports module, linking module, ETL or other tools, etc.). These systems may include

any types of monitors and input devices (e.g., keyboard, mouse, voice recognition, etc.) to enter and/or view information.

[0047] It is to be understood that the software (e.g., auto analytics module, data store services module, fusion module, reporting module, visualization and interactive analysis module, interactive text analytics module, entity resolution module, etc.) of the present invention embodiments may be implemented in any desired computer language and could be developed by one of ordinary skill in the computer arts based on the functional descriptions contained in the specification and flow diagrams or charts illustrated in the drawings. Further, any references herein of software performing various functions generally refer to computer systems or processors performing those functions under software control. The computer systems of the present invention embodiments may alternatively be implemented by any type of hardware and/or other processing circuitry.

[0048] The various functions of the computer or other processing systems may be distributed in any manner among any number of software and/or hardware modules or units, processing or computer systems and/or circuitry, where the computer or processing systems may be disposed locally or remotely of each other and communicate via any suitable communications medium (e.g., LAN, WAN, Intranet, Internet, hardwire, modem connection, wireless, etc.). For example, the functions of the present invention embodiments may be distributed in any manner among the various end-user/client, source, target, and server systems, and/or any other intermediary processing devices. The software and/or algorithms described above and illustrated in the flow diagram or chart may be modified in any manner that accomplishes the functions described herein. In addition, the functions in the flow diagram or chart or description may be performed in any order that accomplishes a desired operation.

[0049] The software of the present invention embodiments (e.g., auto analytics module, data store services module, fusion module, reporting module, visualization and interactive analysis module, interactive text analytics module, entity resolution module, etc.) may be available on a non-transitory computer useable medium (e.g., magnetic or optical mediums, magneto-optic mediums, floppy diskettes, CD-ROM, DVD, memory devices, etc.) of a stationary or portable program product apparatus or device for use with stand-alone systems or systems connected by a network or other communications medium.

[0050] The communication network linking any two or more computing devices may be implemented by any number of any type of communications network (e.g., LAN, WAN, Internet, Intranet, VPN, etc.). The computer or other processing systems of the present invention embodiments may include any conventional or other communications devices to communicate over the network via any conventional or other protocols. The computer or other processing systems may utilize any type of connection (e.g., wired, wireless, etc.) for access to the network. Local communication media may be implemented by any suitable communication media (e.g., local area network (LAN), hardwire, wireless link, Intranet, etc.).

[0051] The system may employ any number of any conventional or other databases or storage structures (e.g., files, databases, data structures, data or other repositories, etc.) to store information (e.g., source data, analytic data, data models, structured and unstructured data, etc.). The data store may

be implemented by any number of any conventional or other databases, or storage structures (e.g., files, databases, data structures, data or other repositories, etc.) to store information (e.g., source data, analytic data, data models, structured and unstructured data, etc.). The data store may be included within or coupled to a server, source, target, and/or client systems. The data store may be remote from or local to the computer or other processing systems, and may store any desired data (e.g., source data, analytic data, data models, structured and unstructured data, etc.).

[0052] The source and analytic data may be arranged within the data store in any desired manner that facilitates suitable searching, analysis and other processing of the data (e.g., extracting entity and relationship data). For example, the data store may include any quantity of areas to store, process, and/or migrate the data. The areas may each be of any quantity, include any suitable storage capacity, store any desired data, and employ any data models (e.g., source, target, or other data model, etc.). The analytic data may include any suitable types of facets comprising information that links data items or instances of data based upon identified relationships between the instances of data.

[0053] The present invention embodiments may employ any number of any type of user interface (e.g., Graphical User Interface (GUI), command-line, prompt, etc.) for obtaining or providing information (e.g., information associated with summary reports generated by the reporting module, etc.), where the interface may include any information arranged in any fashion suitable for enabling interaction with an analyst and filtering of data. The interface may include any number of any types of input or actuation mechanisms (e.g., buttons, icons, fields, boxes, links, etc.) disposed at any locations to enter/display information and initiate desired actions via any suitable input devices (e.g., mouse, keyboard, etc.). The interface screens may include any suitable actuators (e.g., links, tabs, etc.) to navigate between the screens in any fashion.

[0054] The summary reports generated by the reporting module may be visualized within user interfaces to include any information arranged in any fashion, and may be configurable based on rules or other criteria to provide desired information to a user (e.g., based upon filtering or other characteristic information input by the analyst, etc.).

[0055] The present invention embodiments are not limited to the specific tasks or algorithms described above, but may be utilized for reducing a large amount of data to a size available for interactive analysis for any suitable purpose.

[0056] The terminology used herein is for the purpose of describing particular embodiments only and is not intended to be limiting of the invention. As used herein, the singular forms “a”, “an” and “the” are intended to include the plural forms as well, unless the context clearly indicates otherwise. It will be further understood that the terms “comprises”, “comprising”, “includes”, “including”, “has”, “have”, “having”, “with” and the like, when used in this specification, specify the presence of stated features, integers, steps, operations, elements, and/or components, but do not preclude the presence or addition of one or more other features, integers, steps, operations, elements, components, and/or groups thereof.

[0057] The corresponding structures, materials, acts, and equivalents of all means or step plus function elements in the claims below are intended to include any structure, material, or act for performing the function in combination with other claimed elements as specifically claimed. The description of

the present invention has been presented for purposes of illustration and description, but is not intended to be exhaustive or limited to the invention in the form disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the invention. The embodiment was chosen and described in order to best explain the principles of the invention and the practical application, and to enable others of ordinary skill in the art to understand the invention for various embodiments with various modifications as are suited to the particular use contemplated.

[0058] The descriptions of the various embodiments of the present invention have been presented for purposes of illustration, but are not intended to be exhaustive or limited to the embodiments disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the described embodiments. The terminology used herein was chosen to best explain the principles of the embodiments, the practical application or technical improvement over technologies found in the marketplace, or to enable others of ordinary skill in the art to understand the embodiments disclosed herein.

[0059] The present invention may be a system, a method, and/or a computer program product. The computer program product may include a computer readable storage medium (or one or more computer readable storage media) having computer readable storage code or program instructions thereon for causing a processor to carry out aspects of the present invention.

[0060] The computer readable storage medium can be a tangible device that can retain and store instructions for use by an instruction execution device. The computer readable storage medium may be, for example, but is not limited to, an electronic storage device, a magnetic storage device, an optical storage device, an electromagnetic storage device, a semiconductor storage device, or any suitable combination of the foregoing. A non-exhaustive list of more specific examples of the computer readable storage medium includes the following: a portable computer diskette, a hard disk, a random access memory (RAM), a read-only memory (ROM), an erasable programmable read-only memory (EPROM or Flash memory), a static random access memory (SRAM), a portable compact disc read-only memory (CD-ROM), a digital versatile disk (DVD), a memory stick, a floppy disk, a mechanically encoded device such as punch-cards or raised structures in a groove having instructions recorded thereon, and any suitable combination of the foregoing. A computer readable storage medium, as used herein, is not to be construed as being transitory signals per se, such as radio waves or other freely propagating electromagnetic waves, electromagnetic waves propagating through a waveguide or other transmission media (e.g., light pulses passing through a fiber-optic cable), or electrical signals transmitted through a wire.

[0061] Computer readable program instructions described herein can be downloaded to respective computing/processing devices from a computer readable storage medium or to an external computer or external storage device via a network, for example, the Internet, a local area network, a wide area network and/or a wireless network. The network may comprise copper transmission cables, optical transmission fibers, wireless transmission, routers, firewalls, switches, gateway computers and/or edge servers. A network adapter card or network interface in each computing/processing device receives computer readable program instructions from the

network and forwards the computer readable program instructions for storage in a computer readable storage medium within the respective computing/processing device.

[0062] Computer readable program instructions for carrying out operations of the present invention may be assembler instructions, instruction-set-architecture (ISA) instructions, machine instructions, machine dependent instructions, microcode, firmware instructions, state-setting data, or either source code or object code written in any combination of one or more programming languages, including an object oriented programming language such as Smalltalk, C++ or the like, and conventional procedural programming languages, such as the “C” programming language or similar programming languages. The computer readable program instructions may execute entirely on the user's computer, partly on the user's computer, as a stand-alone software package, partly on the user's computer and partly on a remote computer or entirely on the remote computer or server. In the latter scenario, the remote computer may be connected to the user's computer through any type of network, including a local area network (LAN) or a wide area network (WAN), or the connection may be made to an external computer (for example, through the Internet using an Internet Service Provider). In some embodiments, electronic circuitry including, for example, programmable logic circuitry, field-programmable gate arrays (FPGA), or programmable logic arrays (PLA) may execute the computer readable program instructions by utilizing state information of the computer readable program instructions to personalize the electronic circuitry, in order to perform aspects of the present invention.

[0063] Aspects of the present invention are described herein with reference to one or more flowchart illustrations and/or block diagrams of methods, apparatus (systems), and computer program products according to embodiments of the invention. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, can be implemented by computer readable program instructions.

[0064] These computer readable program instructions may be provided to a processor of a general purpose computer, special purpose computer, or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor of the computer or other programmable data processing apparatus, create means for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks. These computer readable program instructions may also be stored in a computer readable storage medium that can direct a computer, a programmable data processing apparatus, and/or other devices to function in a particular manner, such that the computer readable storage medium having instructions stored therein comprises an article of manufacture including instructions which implement aspects of the function/act specified in the flowchart and/or block diagram block or blocks.

[0065] The computer readable program instructions may also be loaded onto a computer, other programmable data processing apparatus, or other device to cause a series of operational steps to be performed on the computer, other programmable apparatus or other device to produce a computer implemented process, such that the instructions which execute on the computer, other programmable apparatus, or other device implement the functions/acts specified in the flowchart and/or block diagram block or blocks.

[0066] The flowchart and block diagrams in the drawings illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the present invention. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of instructions, which comprises one or more executable instructions for implementing the specified logical function(s). In some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts or carry out combinations of special purpose hardware and computer instructions.

What is claimed:

1. A computer-implemented method of processing a request for data comprising:
 - identifying data within a data store satisfying the request, wherein an amount of the identified data exceeds a threshold;
 - analyzing the identified data to produce an initial result set with a reduced amount of data including information summarizing the identified data;
 - generating one or more succeeding result sets each from a corresponding prior result set until a succeeding result set includes an amount of data satisfying the threshold, wherein each succeeding result set includes a reduced

- amount of data relative to the corresponding prior result set, and generating each succeeding result set includes: filtering the corresponding prior result set and generating the succeeding result set including information summarizing the filtered result set.
- 2. The computer-implemented method of claim 1, further comprising:
 - providing the identified data associated with the information from the succeeding result set satisfying the threshold.
- 3. The computer-implemented method of claim 1, wherein filtering the corresponding prior result set includes: filtering the prior result set based on one or more criteria from a user.
- 4. The computer-implemented method of claim 1, wherein the threshold indicates an amount of data enabling visual analysis of the identified data.
- 5. The computer-implemented method of claim 1, wherein the data store includes structured and unstructured data.
- 6. The computer-implemented method of claim 1, wherein the data store includes unstructured data, and identifying data within the data store satisfying the request includes:
 - extracting entity and relationship data from the unstructured data, wherein the information summarizing one or more from a group of the identified data and the filtered result set is based on the extracted entity and relationship data.
- 7. The computer-implemented method of claim 1, further comprising:
 - displaying each result set within a user interface accessible by a user, wherein the user interface facilitates depiction of one or more result sets in a plurality of different formats.

* * * * *